

TreeLine: An Update-In-Place Key-Value Store for Modern Storage

저자 및 학회

2024. 01. 16

Presentation by Yejin Oh, Jisoo Lee, Zhu Yongjie

yejino@ Dankook.ac.kr, lkhejj1@gmail.com,

Contents

1. Structure of TreeLine
2. Building a TreeLine Environment
3. Future Experimental Plans

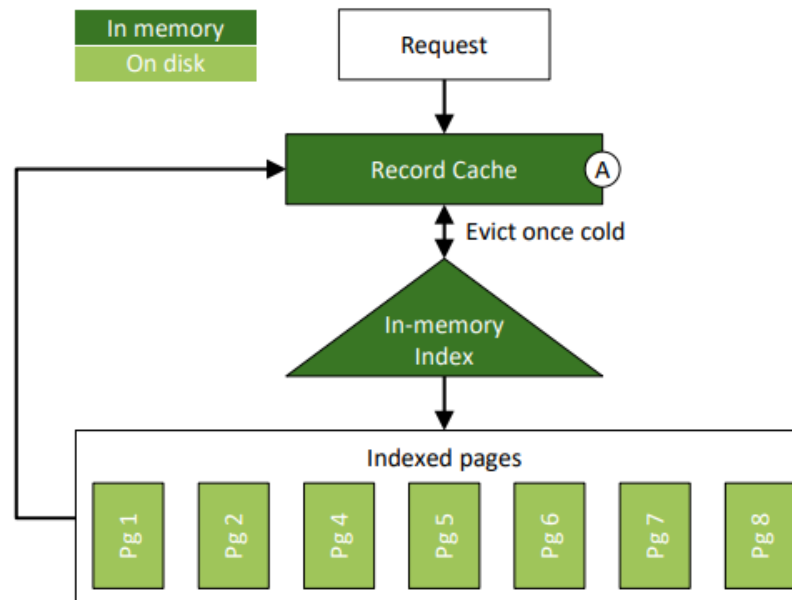
1. Structure of TreeLine

- **What is TreeLine?**
 - An update-in-place key-value store.
 - Well-balanced between read and write performance.
 - Highly efficient in large-scale database systems.

1. Structure of TreeLine

- **Record Caching**

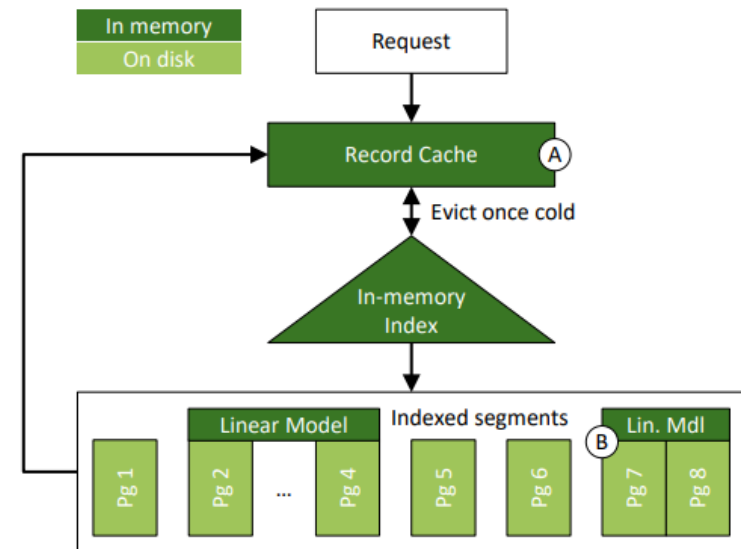

- Optimizes write operations using in-memory caching
- Enhances read performance by maximizing the retention time of hot records in memory



1. Structure of TreeLine

- **Page Grouping**

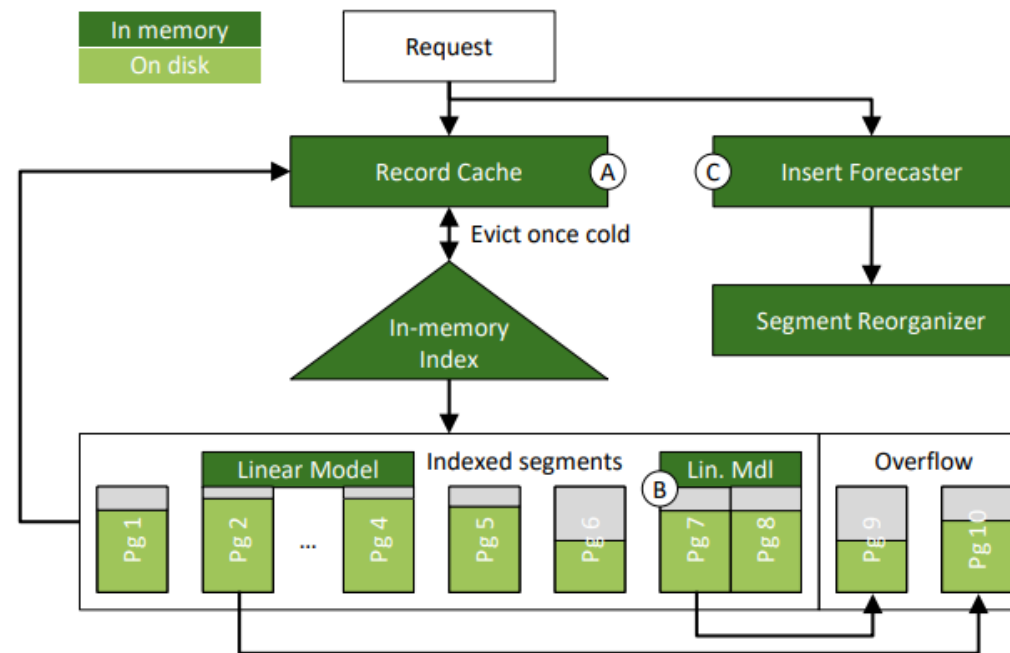
- Groups pages that store adjacent keys
- Stores grouped pages consecutively to improve access efficiency



1. Structure of TreeLine

- **Insert Forecasting**

- Analyzes insertion patterns to predict future insert loads for intelligent space management



2. Building a TreeLine Environment

- TreeLine is an open-source project on GitHub.

- `git clone https://github.com/mitdbg/treeline.git`

- Dependencies

- libtbb-dev
- autoconf
- libjemalloc-dev
- CMake3.17+



`apt install libtbb-dev autoconf libjemalloc-dev`

Create a virtual Python environment using anaconda and "`pip install cmake==3.17`"

2. Building a TreeLine Environment

- TreeLine is an open-source project on GitHub.
 - Compile
 - `mkdir build && build`
 - `cmake -DCMAKE_BUILD_TYPE=Release .. && make -j`
 - `cmake -DCMAKE_BUILD_TYPE=Release -DTL_BUILD_TESTS=ON .. && make -j`
 - `cmake -DCMAKE_BUILD_TYPE=Release -DTL_BUILD_BENCHMARKS=ON .. && make -j`

2. Building a TreeLine Environment

- TreeLine BenchMark

- YCSB

- Go to the " /path/to/db/treeline/build/bench " directory and `"./run_custom"`
 - The workload configuration file is located in " /tree/bench/workload_configs/ "
 - Run `"./run_custom --workload_config= /path/to/db/tree/bench/workload_configs/phased_64B_A_B_A_B_A.yml "`

- Other Bench (adjust options)

- Temporarily **not found**

3. Future Experimental Plans

- How does TreeLine compare against RocksDB and LeanStore on throughput and the amount of physical I/O performed?
- How do the record cache and page grouping contribute to TreeLine's overall performance?
- How does the choice of the page grouping parameters affect the grouping "effectiveness"?
- How effective is insert forecasting?

3. Future Experimental Plans

- YCSB workload

Workload	Description
A	50% Read, 50% Update
B	95% Read, 5% Update
C	100% Read
D	95% Read Latest, 5% Insert
E	95% Range Scan(average length 50, maximum length 100), 5% Insert
F	50% Read-Modify-Write, 50% Read

Thank you