

Updatable Learned Index with Precise Positions Experiment

Wu J, Zhang Y, Chen S, et al., **2021 VLDB**

2024. 02. 28

Presentation by Nakyeong Kim, Suhwan Shin
nkkim@dankook.ac.kr, shshin@dankook.ac.kr

Contents

1. Introduction
2. Previous Experiment Enhancement
(New Observations)
3. Hypothesis
4. Experiments By New Factor
5. Conclusion & Question
6. Future Work

1. Introduction

1) Motivation

Problems with LIPP

- Not tolerate errors
- Create child nodes when conflict occurs (conflict-based structural modification)
- The more conflicts → the higher height of tree → **space amplification**
- Violates the space efficient principle of learned index

Goal

- Analyze the impact of space amplification due to conflicts
- Try to solve it
- + Also, analyze the performance of range query




Figure 1: Trade-off of performance and normalized index size.




Figure 2: Throughput of range queries.

1. Introduction

2) Observe with a focus on Utilization

Most nodes have less than 30% of all entries

→ There exists upper bound

$$T_{\mathcal{M}} = \max_{l \in [0, L-1]} |\{k \in \mathcal{K} | \mathcal{M}(k) == l\}|$$

We observe that there exists an upper bound for the minimum $T_{\mathcal{M}}$, i.e. $\exists \mathcal{M}, T_{\mathcal{M}} \leq \lceil \frac{N}{3} \rceil$ where N is the number of keys in \mathcal{K} , i.e. $N = |\mathcal{K}|$. However, the $\lceil \frac{N}{3} \rceil$ may not be the tightest upper bound in many cases. Thus, our goal is to find a best model $\mathcal{M} = \mathcal{AG}(k) + b$ with the minimum conflict degree $T_{\mathcal{M}}$.

We think that **TM** will be an important factor of space amplification

Node Utilization: Proportion of non-null entries.

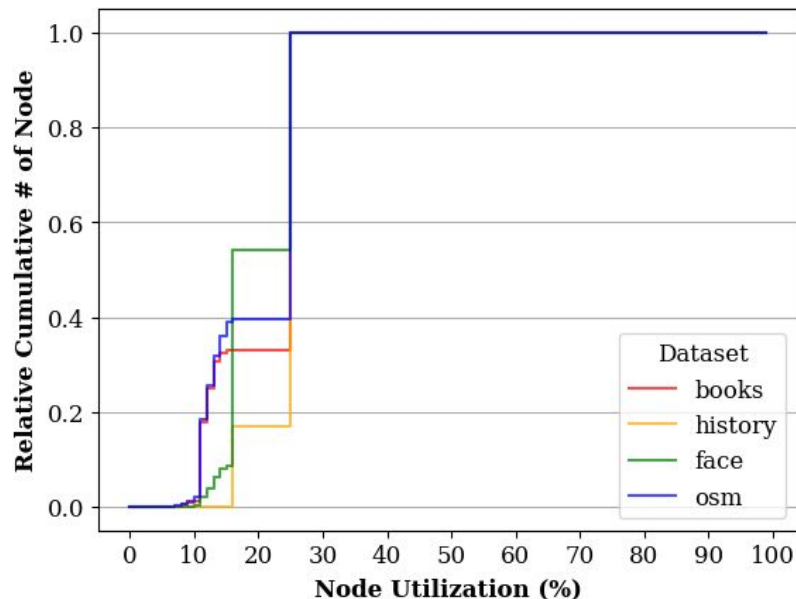


Figure 3: Node utilization CDF

1. Introduction

$$T_{\mathcal{M}} = \max_{l \in [0, L-1]} |\{k \in \mathcal{K} | \mathcal{M}(k) == l\}|$$

3) Parameters of LIPP

- We said that our observation, node utilization upper-bound, is caused by TM, but it is not true
- We found that the factor of affecting to utilization is fill factor(initial node size, gap count)
- We assume that controlling this factor will change performance and index size




Figure 3: Node utilization CDF

1. Introduction

Fill factor: How many fill space when given new keys,
inverse of utilization
e.g., if fill factor is 2, node utilization is 1/2

4) To do Works

1. A sensitivity analysis into updatable learned index structure
 - Node utilization management policy: **fill factor**
 - Model: simple linear regression vs kernelized linear regression
 - Conflict resolving: shifting vs chaining
 - SMO(Structural modification operation): cost-benefit(fanout tree) vs conflict-proportion

1. Introduction

4) To do Works

2. Performance comparison between ALEX and LIPP through size
 - Need understanding of fill factor(parameter) each indexes
3. Which techniques are appropriate when considering performance versus space?
 - Create a new index based on that analysis
 - + Conflict resolving : error-controlled approach (shift-chain hybrid)
 - + Concurrency-friendly : semi-ordered




Figure 4: Expected lookup throughput according to size.

2. Observations

1-1) Node Utilization

Node size set when

- (1) Build tree at first (bulk load)
- (2) Rebuild (adjust)



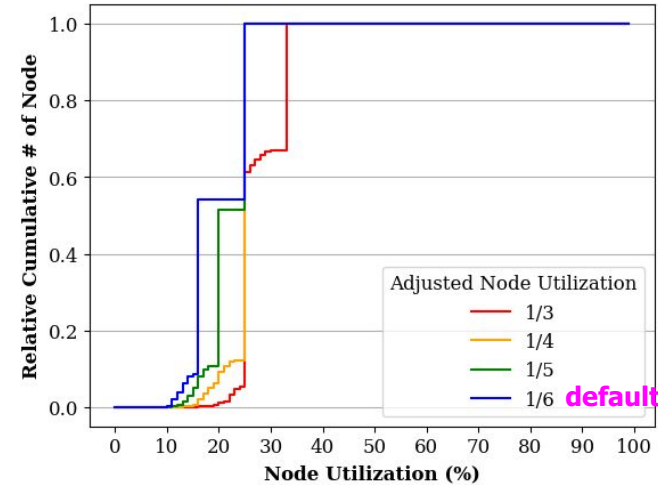



Figure 5: Making “Node” type entry when conflict

2. Observations

Observed after bulk load **100M** keys
Dataset: **Face (1.6GB)**

1-2) Node Utilization CDF

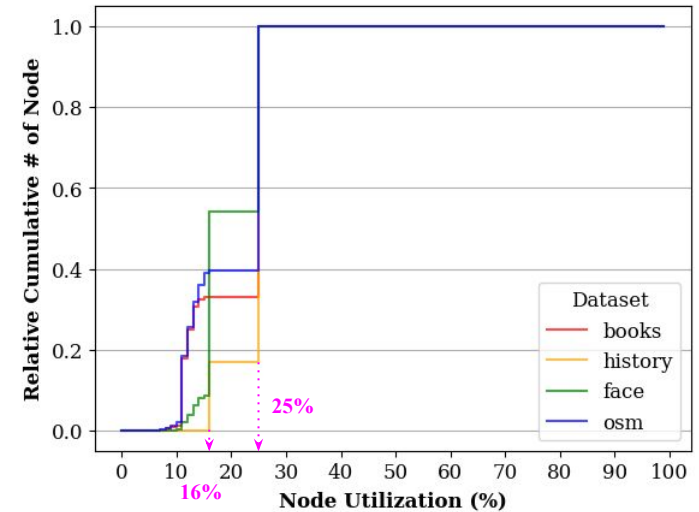


- The initial node size greater, the utilization lower
- The gap size greater, the utilization lower

2. Observations

2) Static Gap Count

- It is important to determine **the array length of new node**, because of trade-off between performance and space consumption
- When new node is first created, gaps are hard-coded with **1,2,5** depending on the keys size
 - 1, size $\geq 1_000_000$
 - 2, size $\geq 100_000$
 - 5, default
- Almost all nodes have a count lower than 100_000, which means that on average, $\frac{1}{6}$ utilized
- It has low hotness



```
const int BUILD_GAP_CNT = compute_gap_count(size);



node->is_two = 0;
node->build_size = size;
node->size = size;
node->fixed = 0;
node->num_inserts = node->num_insert_to_data = 0;

{
    const int L = size * static_cast<int>(BUILD_GAP_CNT + 1);
    node->num_items = L;
}
```

2. Observations

Observed after bulk load **100M** keys
Dataset: **Face (1.6GB)**

3) Root Node Size (Initial Node Size)

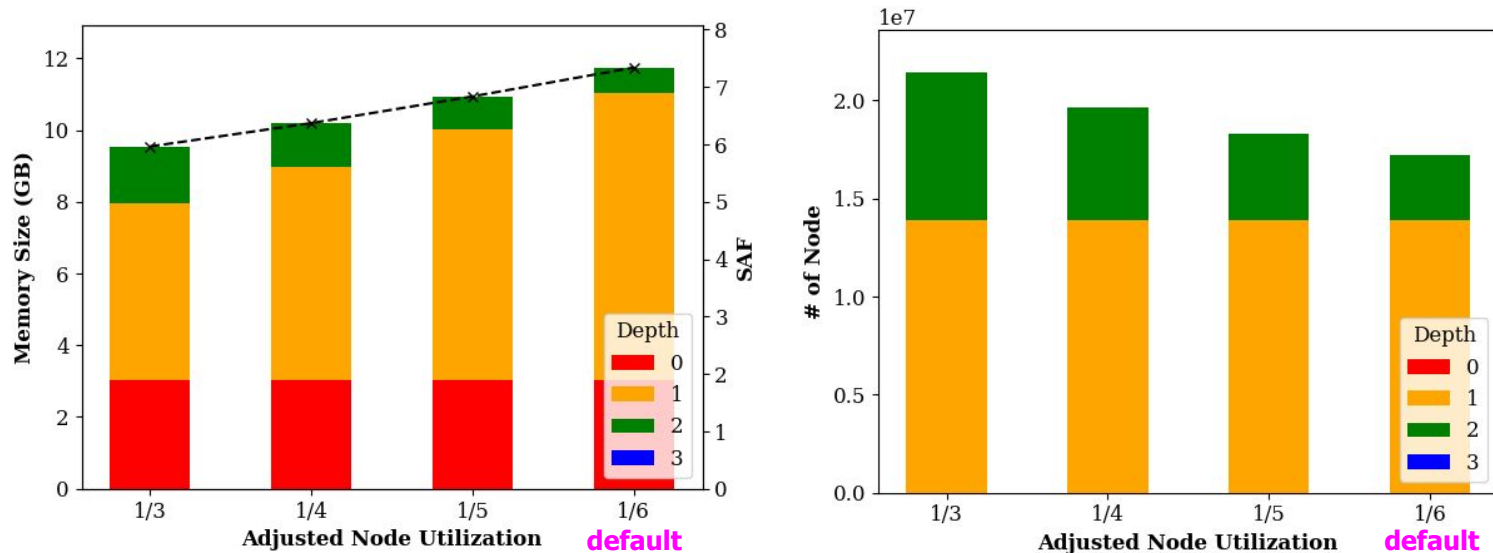


- Root node takes up a significant percentage of the total index size
- Initial node size doesn't have much impact on overall size
- Large number of nodes, but small percentage of size (MBs)

2. Observations

Observed after bulk load **100M** keys
Dataset: **Face**

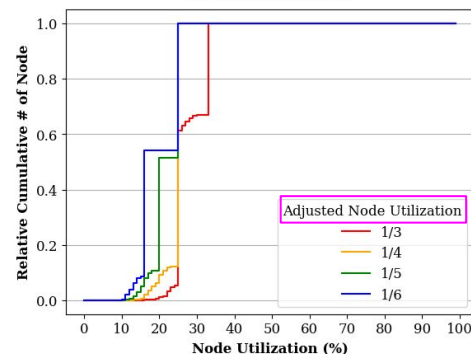
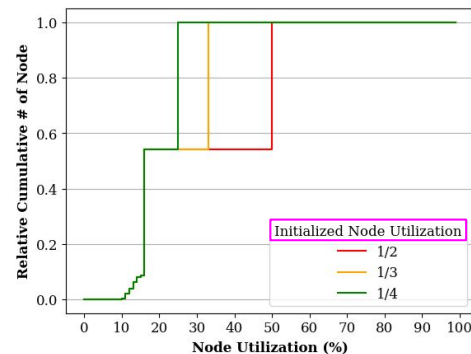

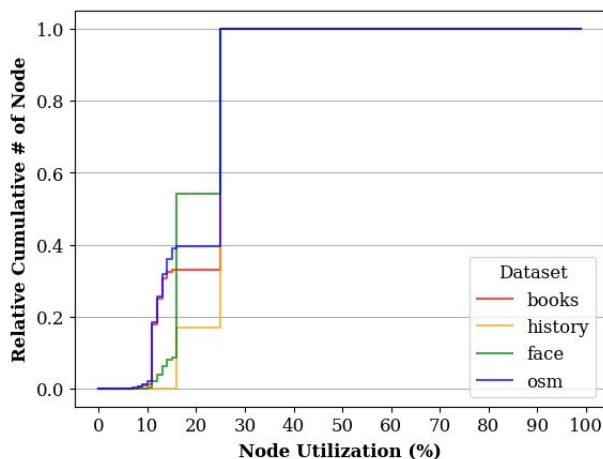
3) Root Node Size (Gap Count)



- Gap count has more impact than initial node size, but not primarily
- There may be some other factors (e.g., α , β , FMCD ...)

3. Hypothesis

- Hypothesis: The cause of LIPP's SAF may be array management policy (node utilization)
- Reason: Node utilization is low
- Verification: Increasing node utilization

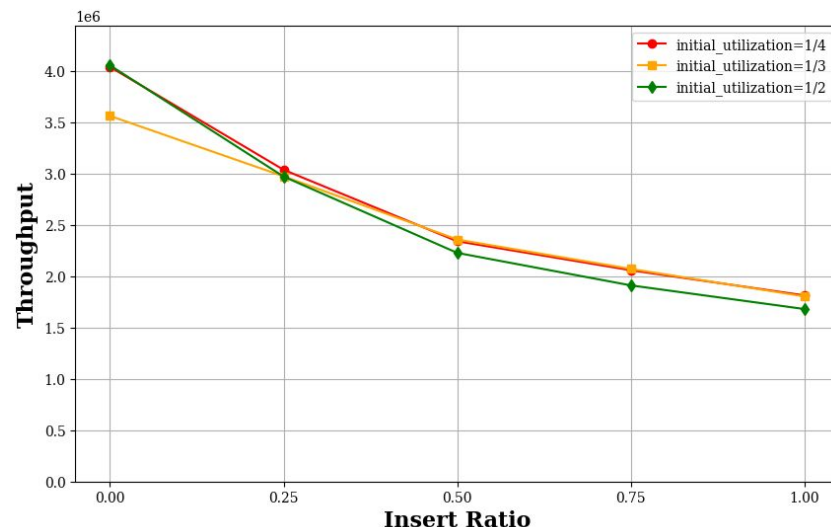
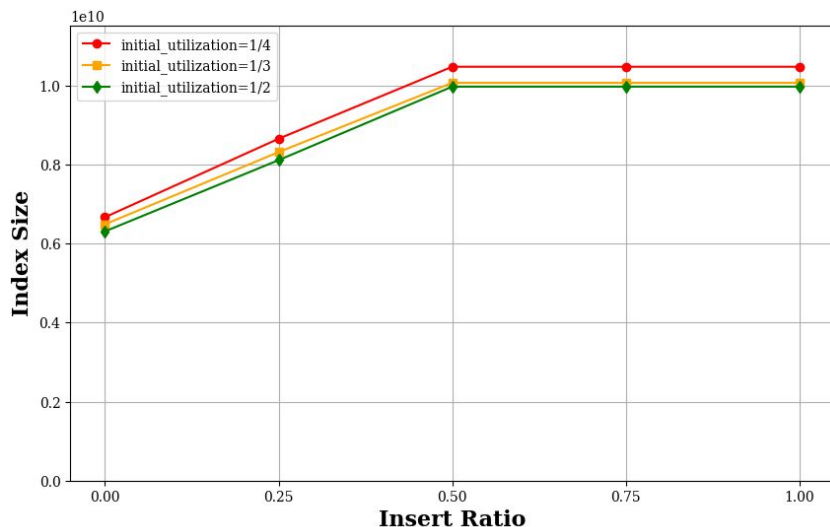


4. Experiments

Observed after bulk load **100M** keys
Dataset: **Face**

Goal : Impact of lipp's array utilization policy (initial node size)

Observation. The **initial node size** has little effect on index size and read-write performance

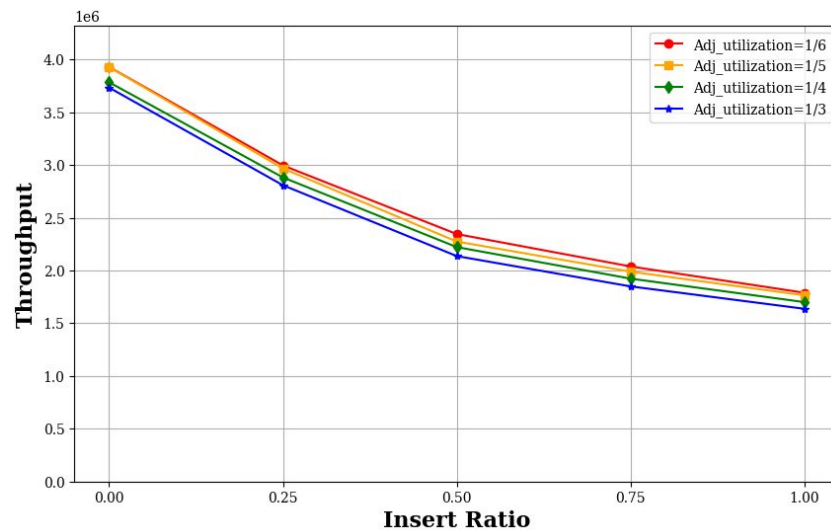
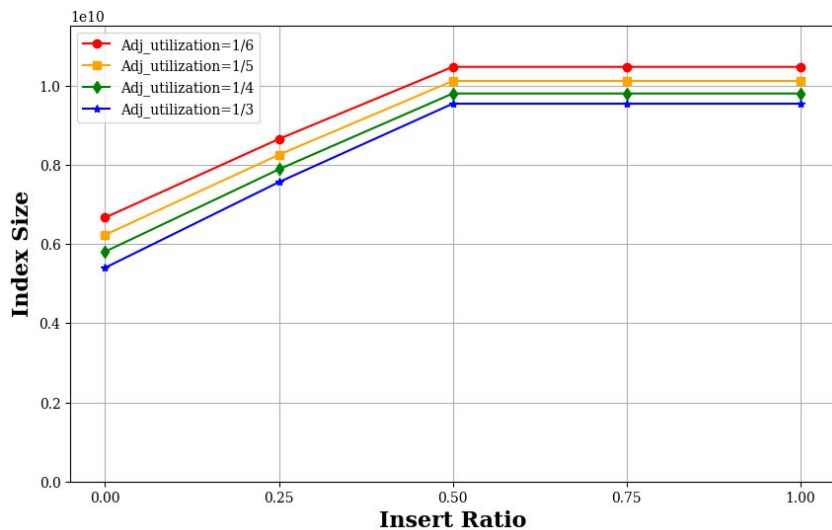


4. Experiments

Observed after bulk load **100M** keys
Dataset: **Face**

Goal : Impact of Lipp's array utilization policy (gap count)

Observation. The **gap count** has a greater impact than the initial node size, but it is not the main cause




5. Conclusion & Question

Experiment conclusion:

- Array management policy was not the main cause of Space Amplification (SAF)

Question: No nodes exceed initial node utilization (25%)

- If fill factor is the problem, $3/8$ should also exist!
- But, node with a utilization rate of $3/8$ is not observed
→ Probably FMCD, conflict issue or something...



6. Future work

1. A sensitivity analysis into updatable learned index structure

- **Current Hypothesis: Array management policy**

- **fill factor: initial node size, gap_count**

- **New Hypothesis (Expectation)**

- **rebuilding condition: α, β**

- **training model: FMCD**

- **Model : simple linear regression vs kernelized linear regression**

- Conflict resolving : shifting vs chaining

- SMO (Structural modification operation) : cost-benefit (fanout tree) vs rebuilding

$$\frac{n.element_num}{n.build_num} \geq \beta \quad \beta \text{ is set to 2 by default}$$

$$\frac{n.conflict_num}{n.element_num - n.build_num} \geq \alpha \quad \text{we set the threshold } \alpha = 0.1$$

2. Performance comparison between ALEX and LIPP through size

- Need understanding of fill factor parameter each indexes

3. Which techniques are appropriate when considering performance versus space?

- Create a new index based on that analysis

- + Conflict resolving : error-controlled approach (shift-chain hybrid)

- + Concurrency-friendly : semi-ordered

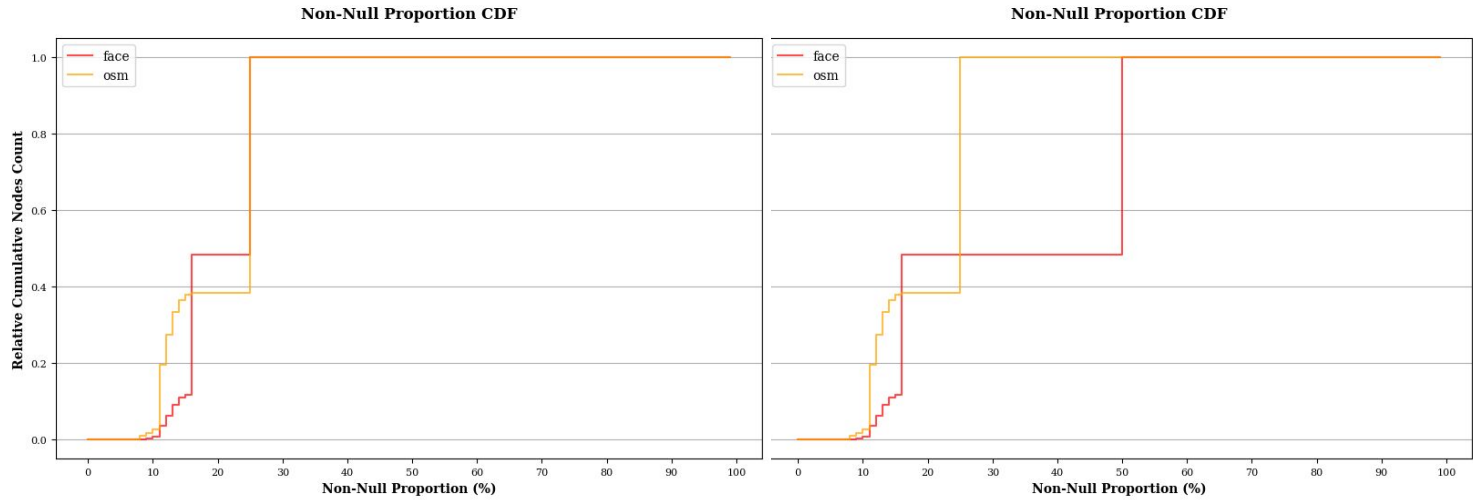
Q&A



Thank you!

2. Observation

3) Ro



4. New? Hypothesis

- Adjustment trigger
 - 1) fill factor
 - 2) conflict num

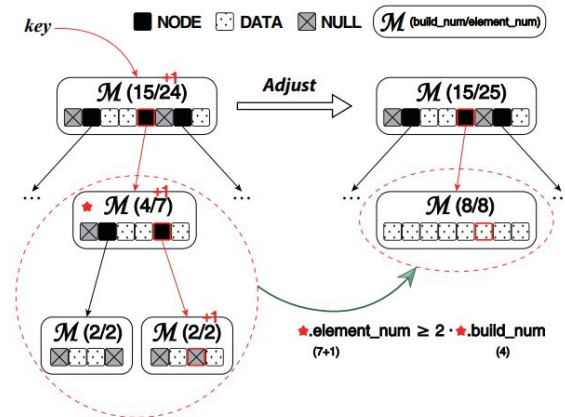


Figure 4: Node Adjustment

- This is because we are considering the utilization of the sub-tree, not the node
- When β is 2, the actual utilization may be much less than 50%

$$\frac{n.\text{element_num}}{n.\text{build_num}} \geq \beta \quad \beta \text{ is set to 2 by default}$$

$$\frac{n.\text{conflict_num}}{n.\text{element_num} - n.\text{build_num}} \geq \alpha \quad \text{we set the threshold } \alpha = 0.1$$

1. Summary

Motivation

Poor understanding of Updatable Learned Index (LIPP)

→ Additional studying: Array Size Policy, Rebuilding Process, FMCD

Analysis of relationship between **index size** and **performance**

→ Between LIPP and ALEX, which structure is more suitable for a learned index?

+ Error-Controlled Approach, Range Query, Semi-Ordered

1.1. Adjustment Strategy

When to Adjust

1. Insert key(s)
2. Update and check statistics of nodes in the traversal path
3. Trigger adjustment on a chosen node when certain conditions are satisfied
 - a. $\frac{n.element_num}{n.build_num} \geq \beta$ β is set to 2 by default
 - b. $\frac{n.conflict_num}{n.element_num - n.build_num} \geq \alpha$ we set the threshold $\alpha = 0.1$

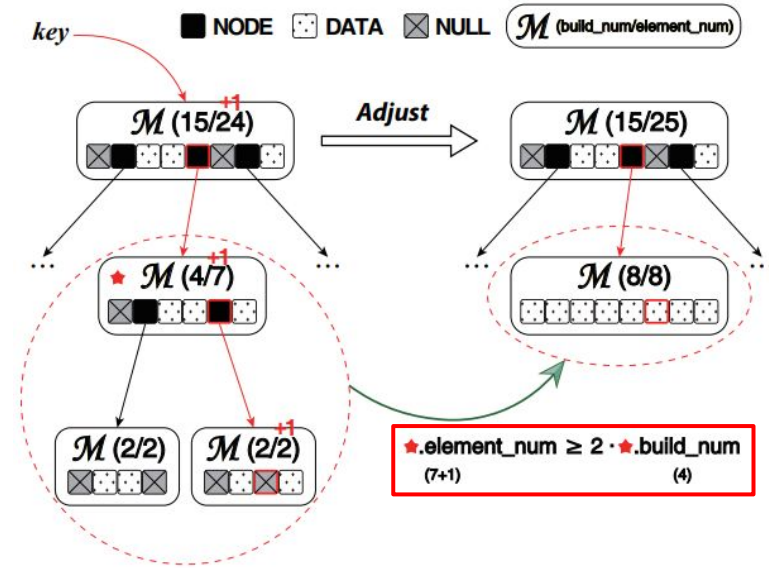


Figure 4: Node Adjustment

1.1. Adjustment Strategy

How to Adjust

1. Collect all elements(keys) in the subtree rooted at node by sequential traversal
2. Build a partial tree on the elements
3. Update the pointer of the original node to point of the new node(tree)

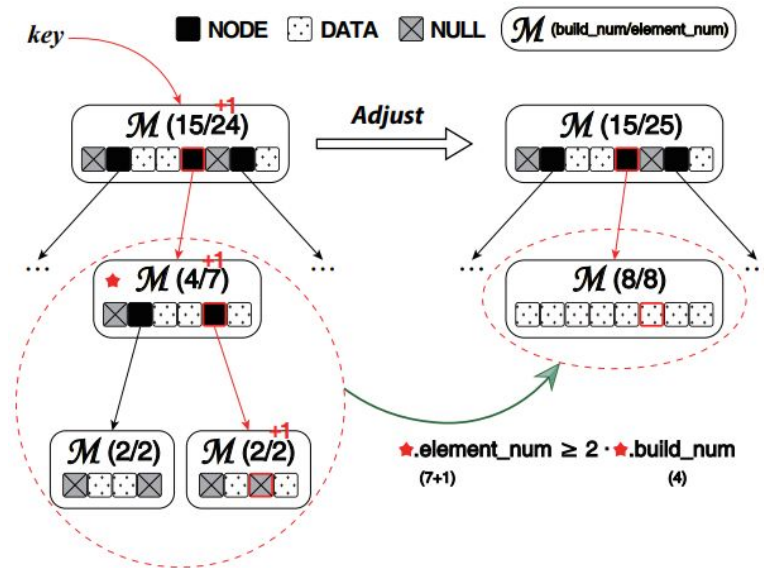


Figure 4: Node Adjustment

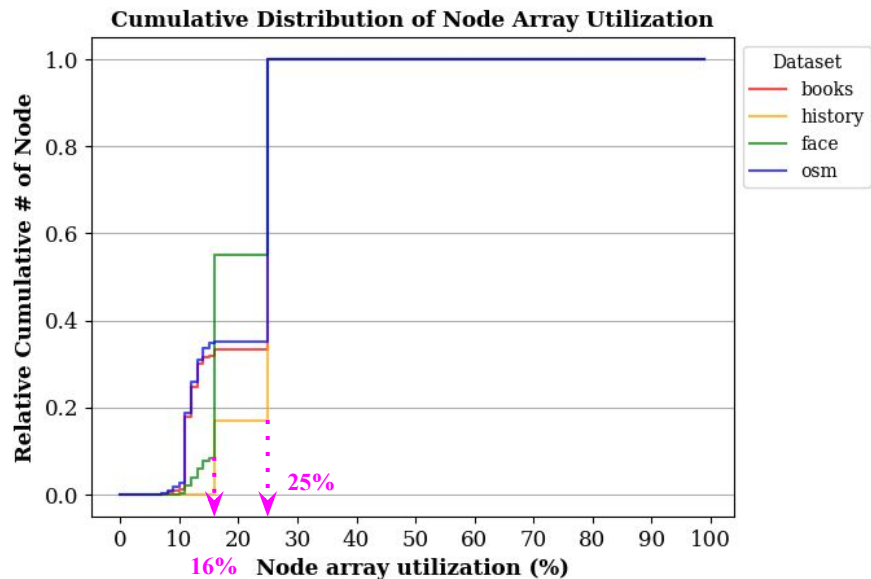

1. Summary

Parameter of Index

Most of raw data is "8,2", which means "total,node" entries.

There are things the paper doesn't show.

- Array size of child node when first created is hard-coded by 8. (0.25 utilization)
- Array size of new node when need rebuilding(or bulk loading) can be 2x, 3x, and 6x, respectively.



1. Summary

Purpose

Lookup/Insert performance comparison according to size

- Tuning parameter: ALEX(fill factor), LIPP(gap count)

LIPP vs ALEX

- Model
- Array Management Policy
- Shift/Chaining
- SMO

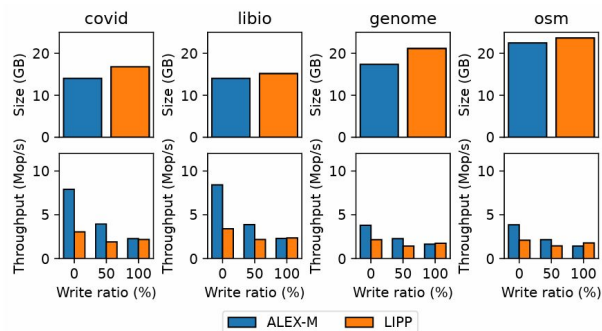




Figure 9: ALEX-M vs. LIPP (when ALEX is tuned to use roughly the same amount of memory as LIPP).

2.1. Experiment


Tendency by Dataset



2.2.1. Experiment

scan_num: 100
iteration: 1M
table_size: 100M (key_range: 200M)
operations: range query only


Range Query Throughput by Dataset



2.2.2. Experiment

scan_num: 100
iteration: 1M
table_size: 100M (key_range: 200M)
operations: range query only







Range Query Latency by Dataset



2.2.2. Experiment

scan_num: 100
iteration: 1M
table_size: 100M (key_range: 200M)
operations: range query only







Lookup Throughput by Dataset



2.2.2. Experiment

Lookup Latency by Dataset


scan_num: 100
iteration: 1M
table_size: 100M (key_range: 200M)
operations: range query only







2.3. Dataset Hardness

New Criteria

- Previous experiments defined data hardness as the number of segments (optimal PLA model)
- But our approach is based on **conflict count**, so it is not compatible
- Need to quantify to other criteria
- Plan to use definition of conflict degree of node





GAP_CNT: 1/2/5





GAP_CNT: 1/2, 1/3, 1/5

GAP_CNT: 0/1/2

num_items: 8, 6, 4

num_items: 8



num_items: 4



2. Observations

Observed after bulk load **100M** keys
Dataset: **Face**

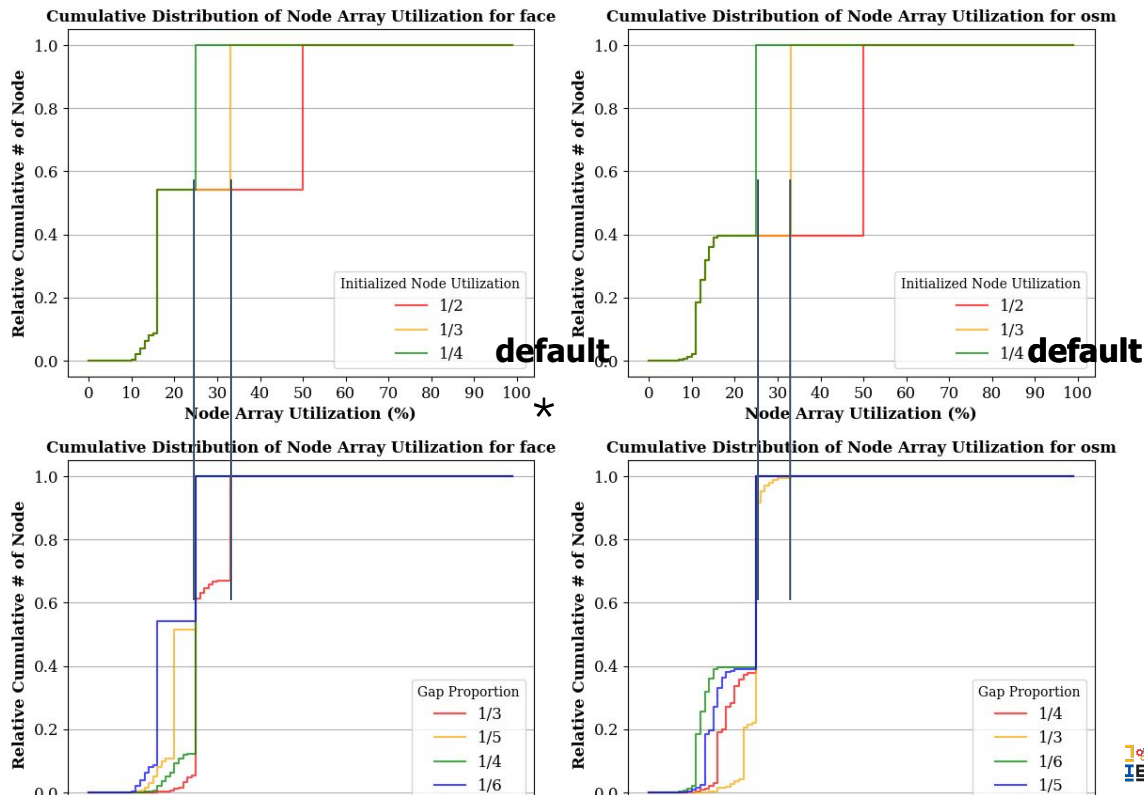
3) Root Node Size (Initial Node Size)



3. Experiment

Observed after bulk load **100M** keys



Array Utilization CDF (Initialized Node Size)



3. Experiment

Observed after bulk load **100M** keys



Array Utilization CDF (Gap Proportion)



3. Experiment

Goal : lipp의 array utilization policy의 영향 (Size)

Observation . 우리는 개 삽질을 한것이다. 알파 베타가 답이었다.





3. Experiment

Goal : lipp의 array utilization policy의 영향 (Performance)

read-write mixed workload



Observation . 모름



3. Experiment

Goal : lipp의 array utilization policy의 영향 (Performance)



Observation . 모름



3. Experiment

Goal : lipp의 array utilization policy의 영향 (Size)

Observation . 모름



Hypothesis

Adjustment trigger

- 1) fill factor
- 2) conflict num

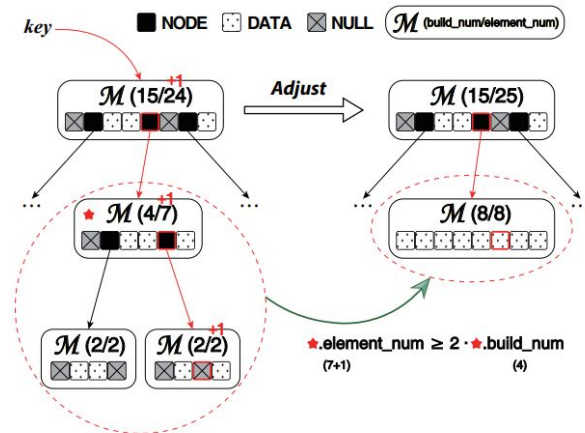



Figure 4: Node Adjustment

해당 노드가 아니라, sub-tree의 utilization을 고려하는 것이므로
beta가 2인 경우의 실제 utilization은 50% 보따 훨씬 작을 수 있음
좀 더 생각해 보는 걸루

fill factor 가 문제면, 3/8도 존재해야함...
하지만 3/8 왜 없누? 누가 범인이누? 몰루?
아마도 FMCD, Conflict 문제인듯

좀 바뀌어야 되긴 할듯요. 정리나 음 새로운 발견?
뭐든. 근데 제가 future work에 쓴거 보시면
가설을 2개로 나눠서 썼는데 보셨어요?

$$\frac{n.\text{element_num}}{n.\text{build_num}} \geq \beta \quad \beta \text{ is set to 2 by default}$$

$$\frac{n.\text{conflict_num}}{n.\text{element_num} - n.\text{build_num}} \geq \alpha \quad \text{we set the threshold } \alpha = 0.1$$

4. Future work

1. A sensitivity analysis into updatable learned index structure
 - **Array management policy : fill factor, initial node size (이 산이 아니다!)**
 - 사실 fill factor는 alpha, beta였던 거임 ㅋㅋ (아님 말고 ㅋ)
 - Model : simple linear regression vs kernelized linear regression (**FMCD**)
 - Conflict resolving : shifting vs chaining
 - SMO (Structural modification operation) : cost-benefit(fanout tree) vs rebuilding
2. Performance comparison between ALEX and LIPP through size
 - Need understanding of fill factor parameter each indexes
3. Which techniques are appropriate when considering performance versus space?
 - Create a new index based on that analysis
 - + Conflict resolving : error-controlled approach (shift-chain hybrid)
 - + Concurrency-friendly : semi-ordered