

FILM: a Fully Learned Index for Larger-than-Memory Databases

Ma, Chaohong, et al. VLDB 2022

2024. 02. 28

Presentation by Yejin Oh, ZhuYongjie, Boseung Kim

yeojinoh@dankook.ac.kr, arashio1111@dankook.ac.kr, bskim1102@dankook.ac.kr

Contents

1. Introduction

2. Design

3. FILM

1) Design of piece

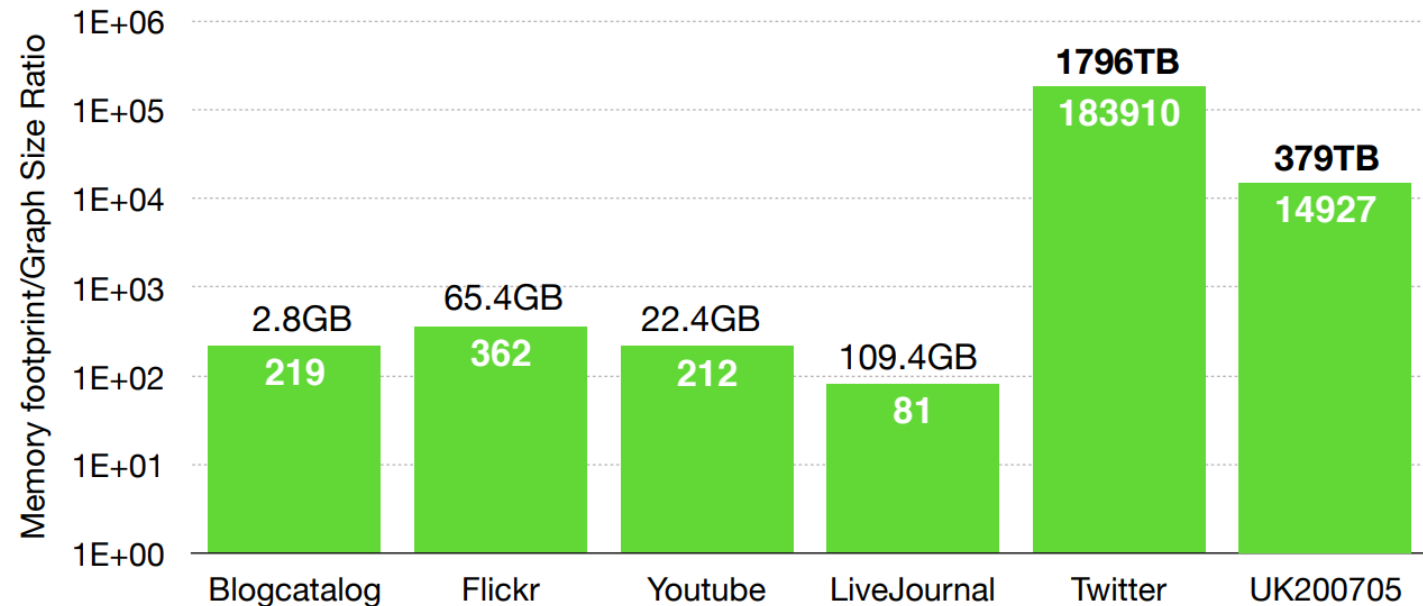
2) Dynamic learned index

3) Adaptive LRU

4. Experiments

Introduction

- Data characteristics
 - Large volume and high velocity
 - Append-only
 - Diverse workloads



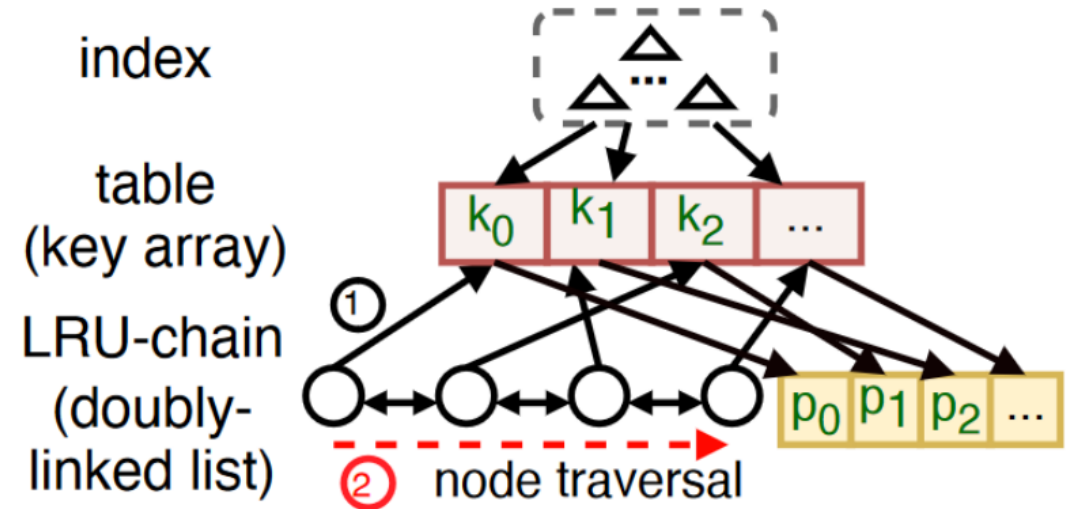
Reference: Shao, Yingxia, et al. "Memory-aware framework for fast and scalable second-order random walk over billion-edge natural graphs." *The VLDB Journal* 30.5 (2021)

Introduction

- Challenge
 - 1) Existing index has **high disk I/O**
 - 2) Existing index occupies **a large portion of memory**
 - 3) In 'anti-caching', computation **overhead in LRU is too high**
 - 4) Learned index assumes data is stored in **a contiguous array**
 - 5) Learned index is **for in-memory** data

Introduction

- Anti-caching
 - To reduce the consumption of main memory
 - High memory consumption
 - High LRU maintenance costs



(a) anti-caching's LRU

Introduction

- Goal
 - Using Learned Indexes in Heterogeneous storage
 - Cost Reduction of Cold Data Identification with Adaptive LRU

Design

1. Lightweight Machine learning model

- Capturing data distribution to reduce memory usage

2. Adaptive LRU

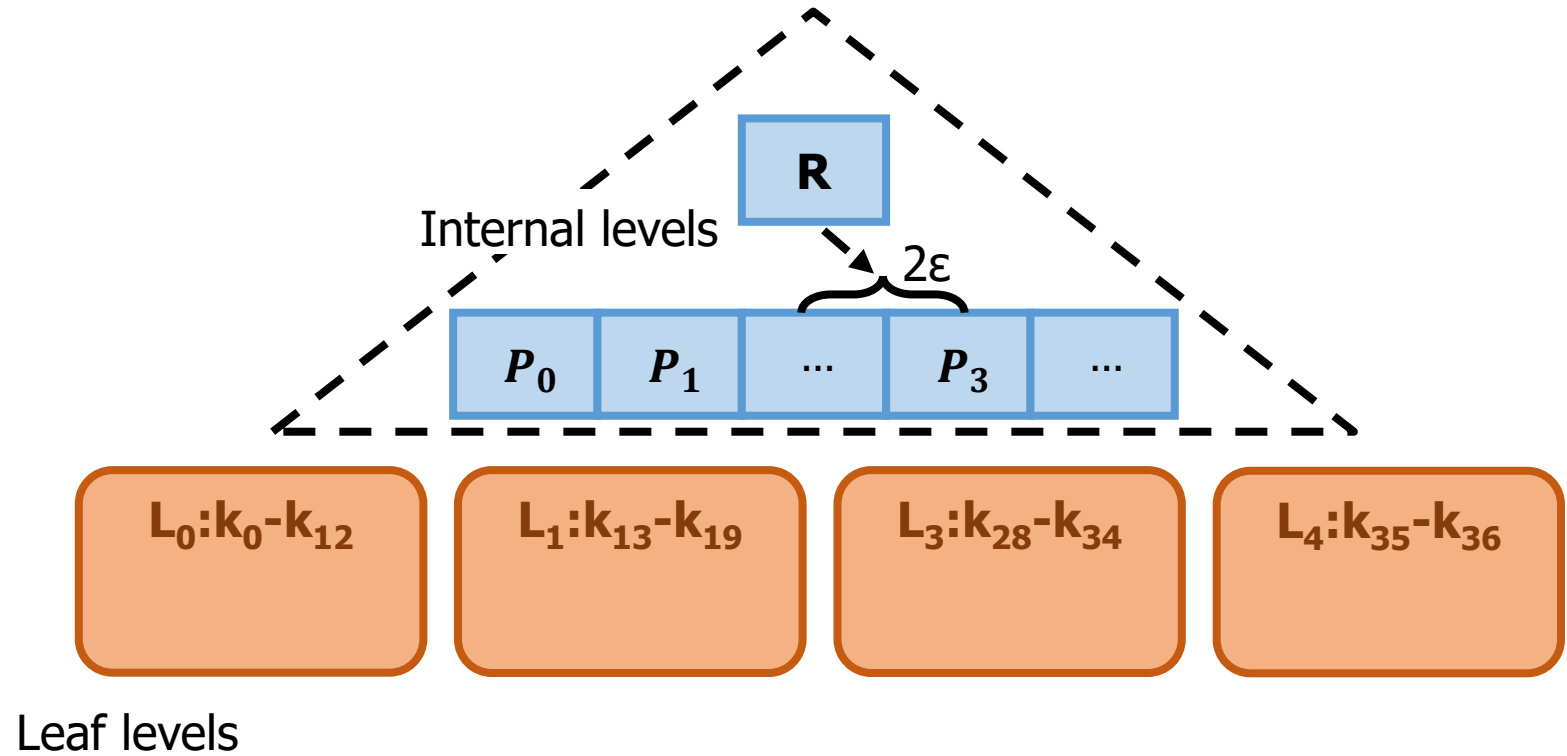
- Adaptive LRU for efficient cold data identification, cost-effective management

3. Unified Tree Structure

- Stable tree structure minimizes maintenance despite data swapping

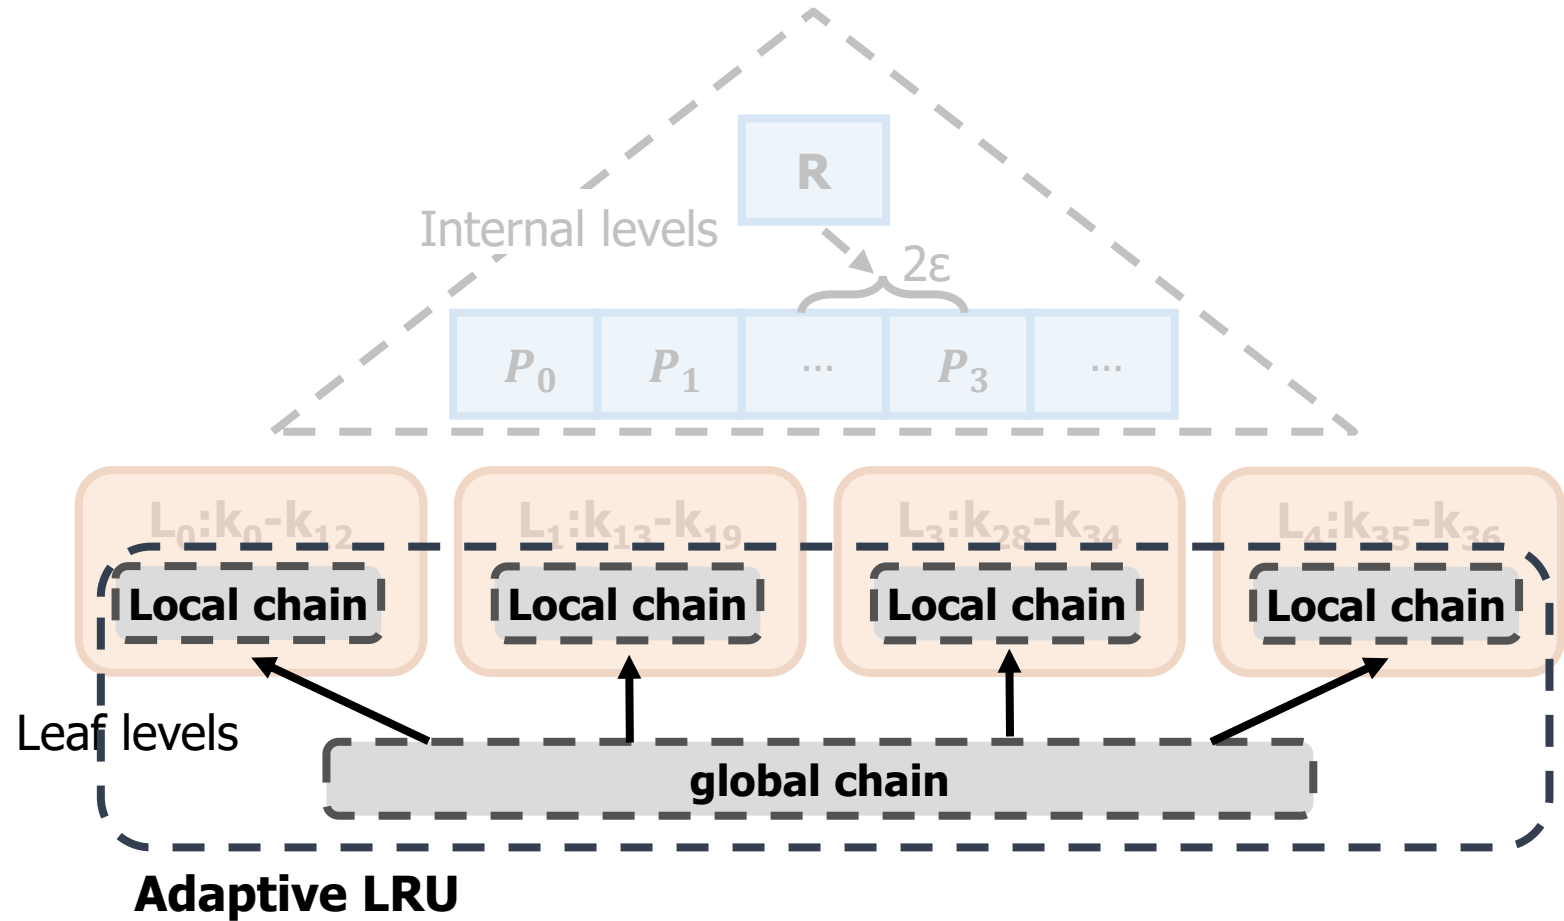
Design

- **Learned index**
- adaptive LRU
- Unified tree structure



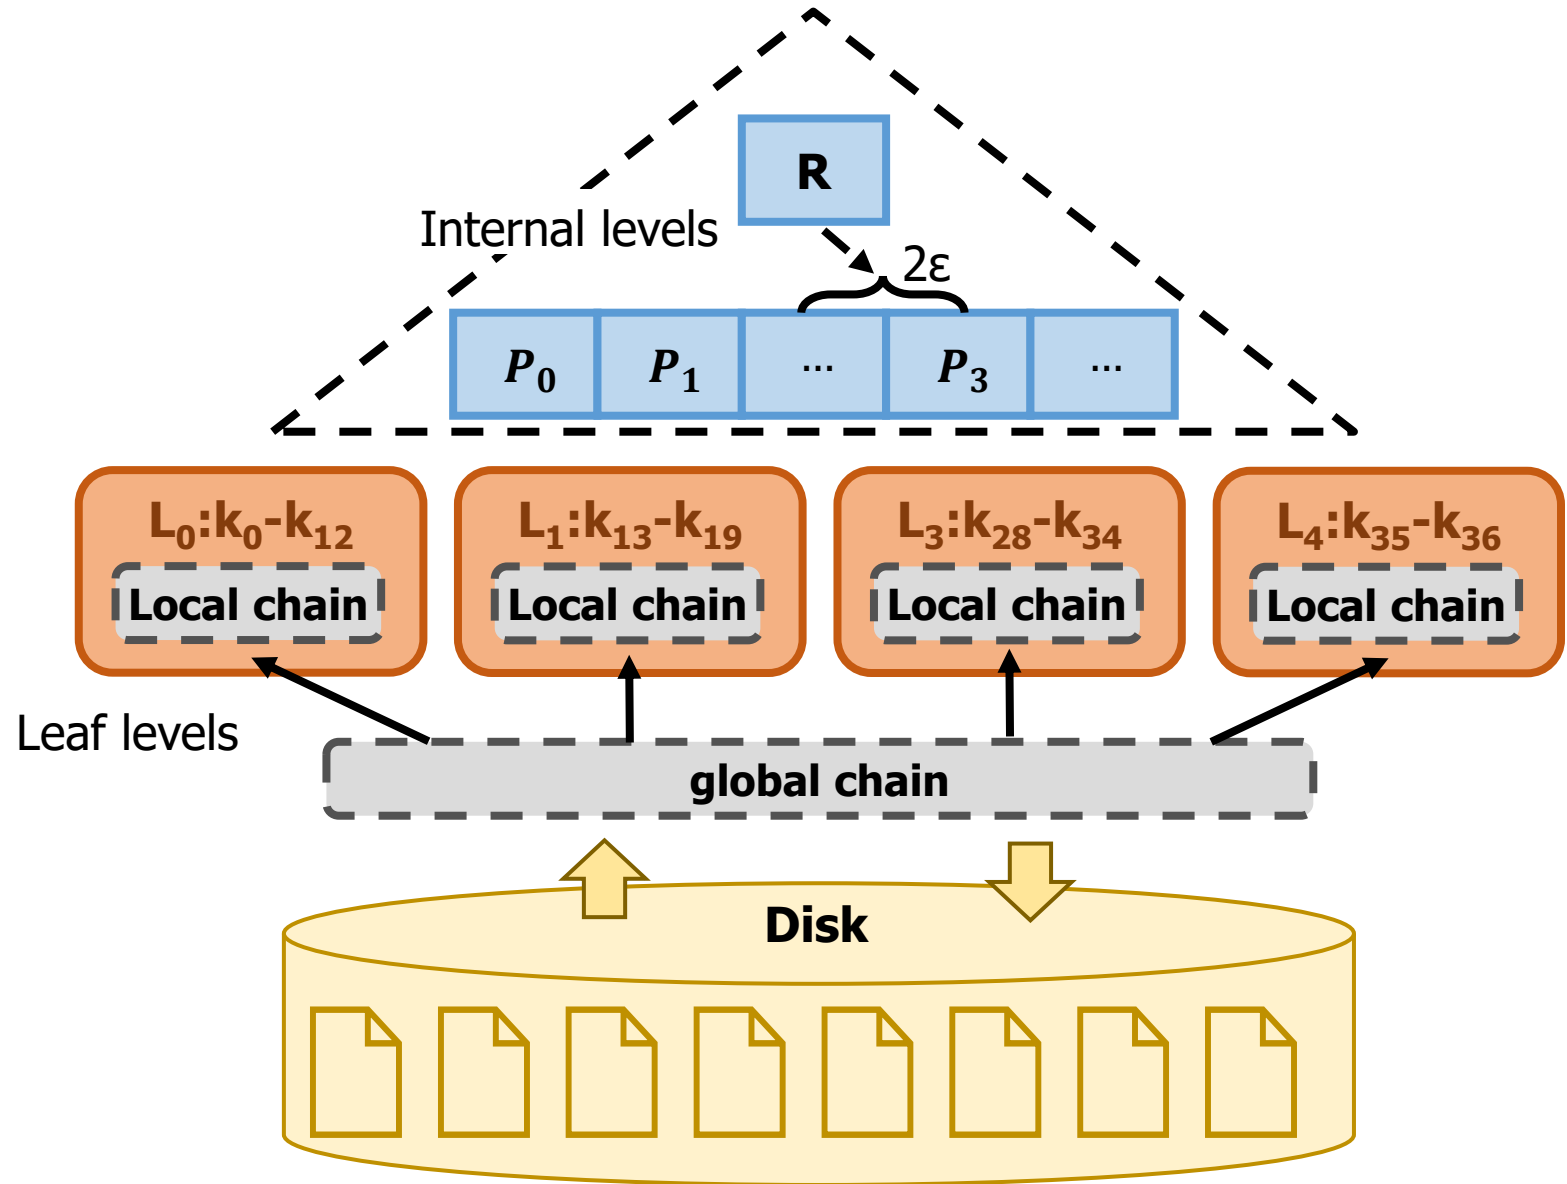
Design

- Learned index
- **adaptive LRU**
- Unified tree structure



Design

- Learned index
- adaptive LRU
- **Unified tree structure**



The Design of Piece

- Piece

- Contains a sub-range of data and an approximation model fitted on the sub-range
- Use a list of pwlfes (piecewise linear functions) to partition the sub-range of keys
- The model guarantees the following formula (ε : a specified distance bound)

$$|pred_pos - true_pos| \leq \varepsilon$$



- Newly inserted key breaking constraint is inserted into new piece, this key is called 'break_k'

The Design of Piece

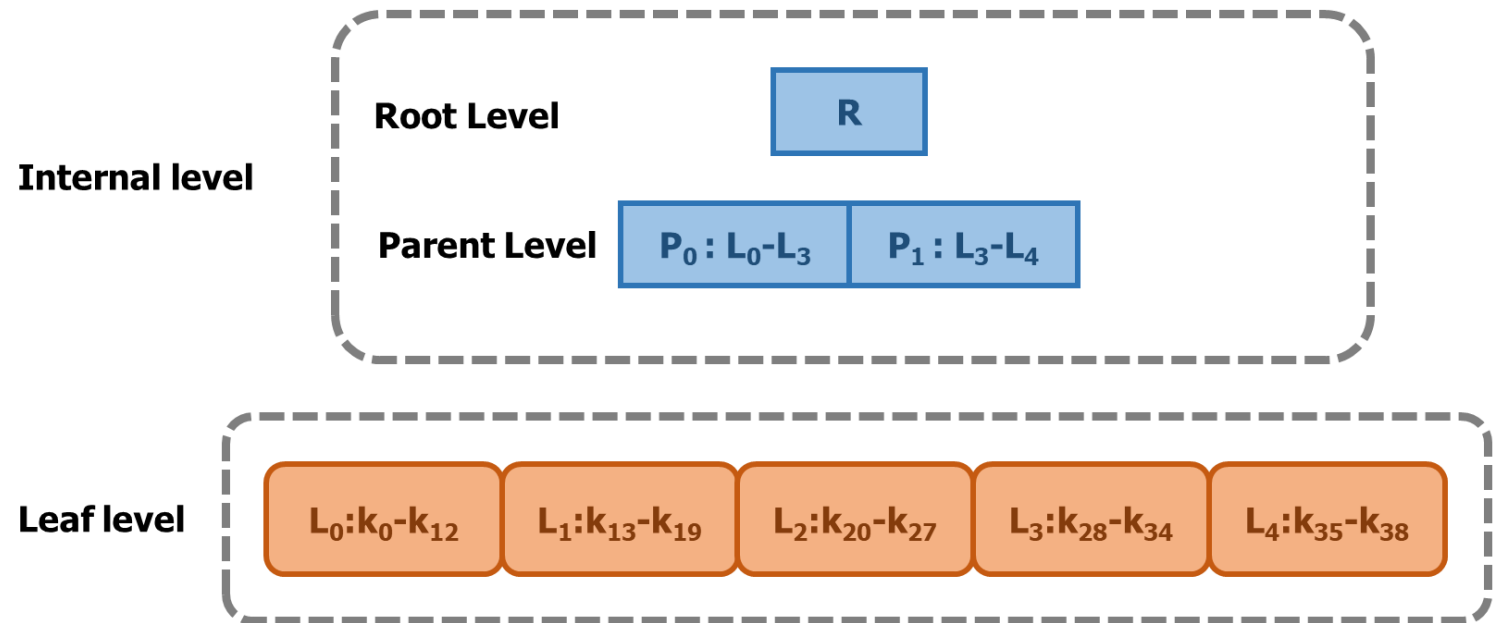
i_piece 

l_piece 

- Two types of piece
 - Component of pieces



- Learned model with pieces

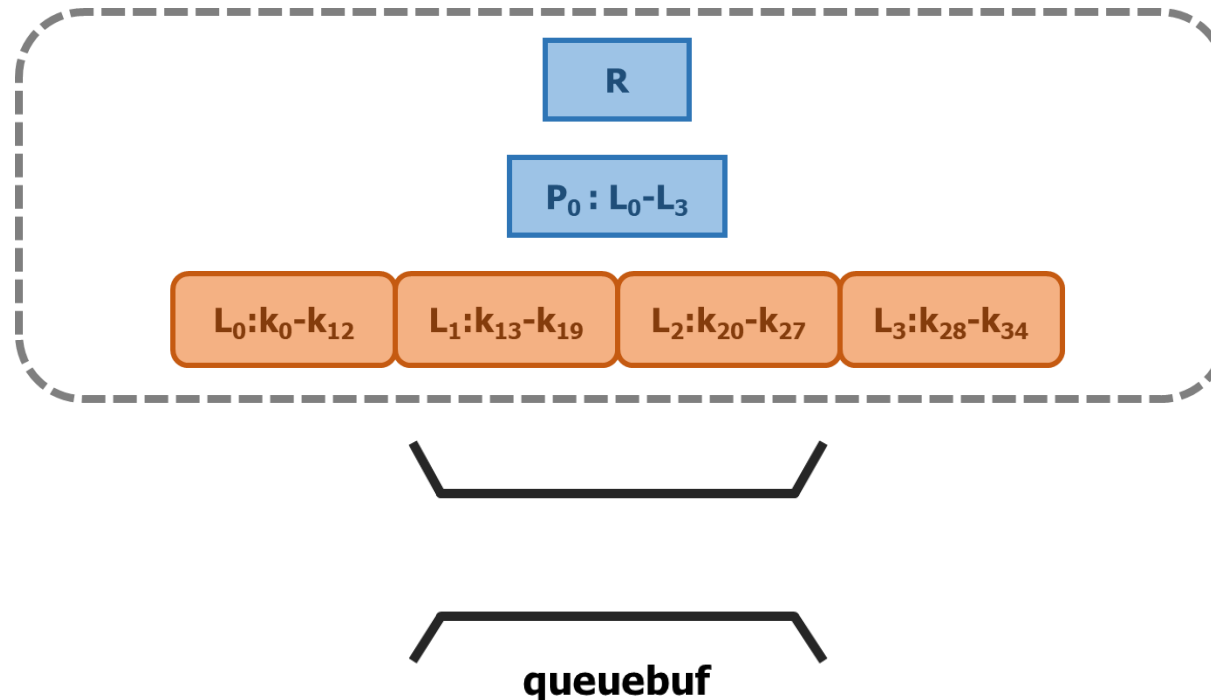


Dynamic Learned Model

- The construction process of the learned model
 - Two break_k are needed to create l_piece
 - Queuebuf is designed to temporarily store the break_k

i_piece 

l_piece 

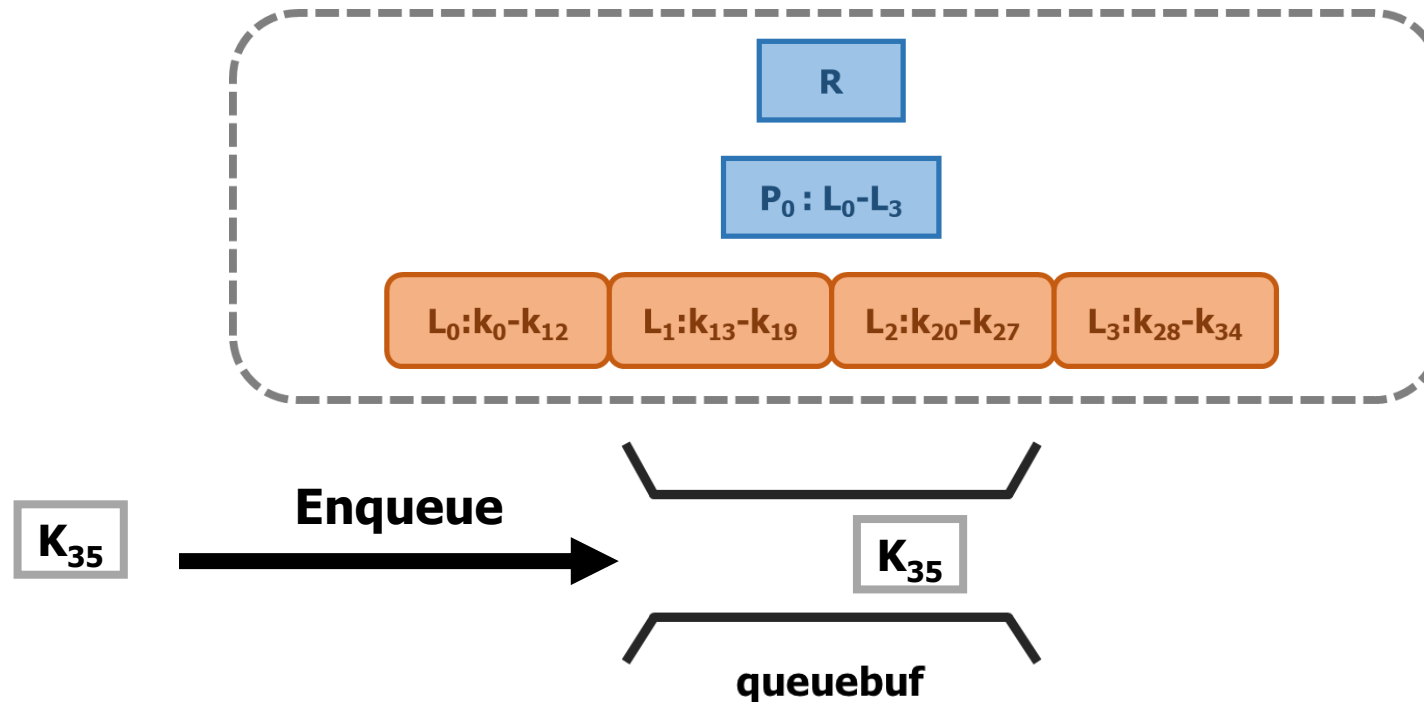


Dynamic Learned Model

- The construction process of the learned model
 - Put newly arrived key (break_k) into queuebuf
 - If Key is not break_k, directly put into l_piece

i_piece 

l_piece 

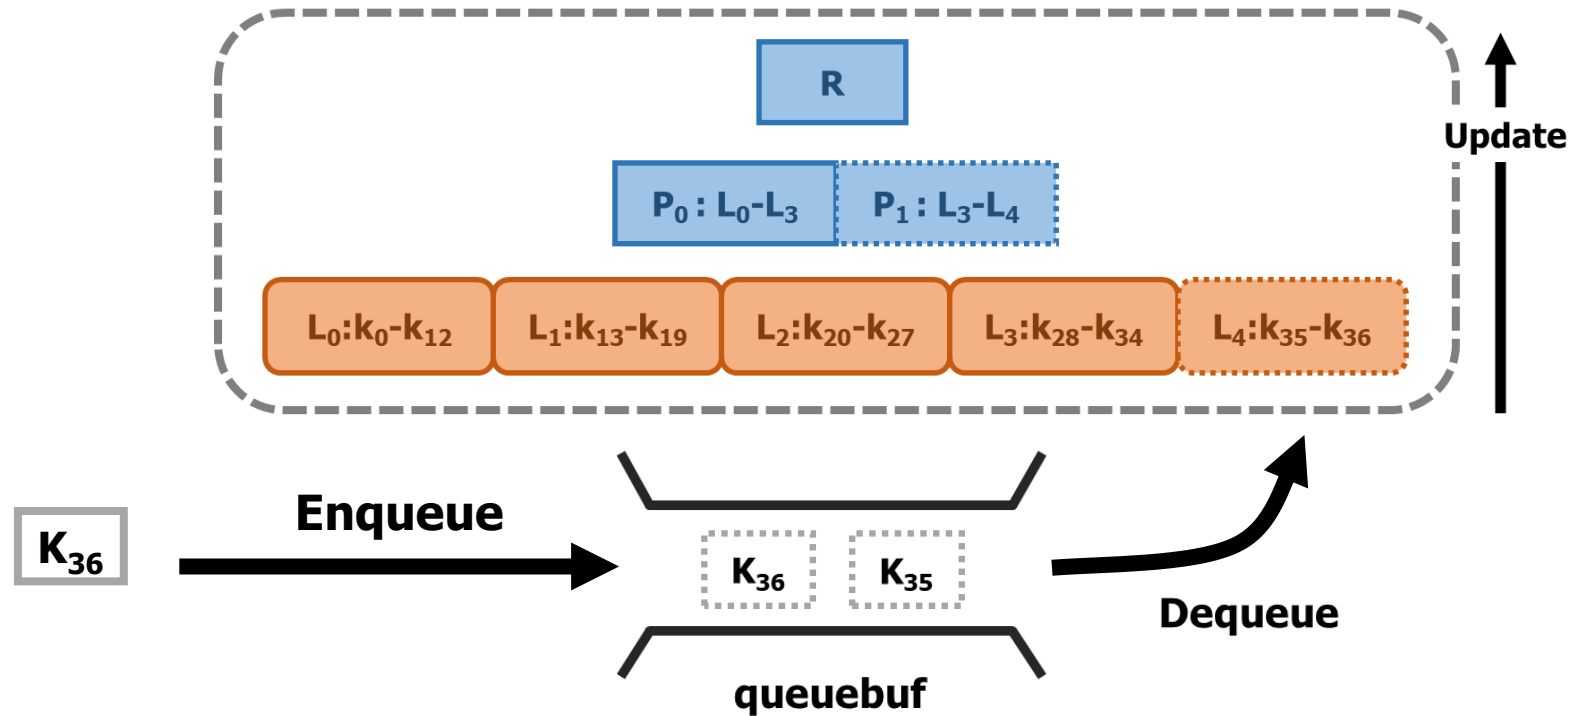


Dynamic Learned Model

- The construction process of the learned model
 - Queuebuf is full, keys are removed and create a I_piece
 - Check if the parent level should be updated

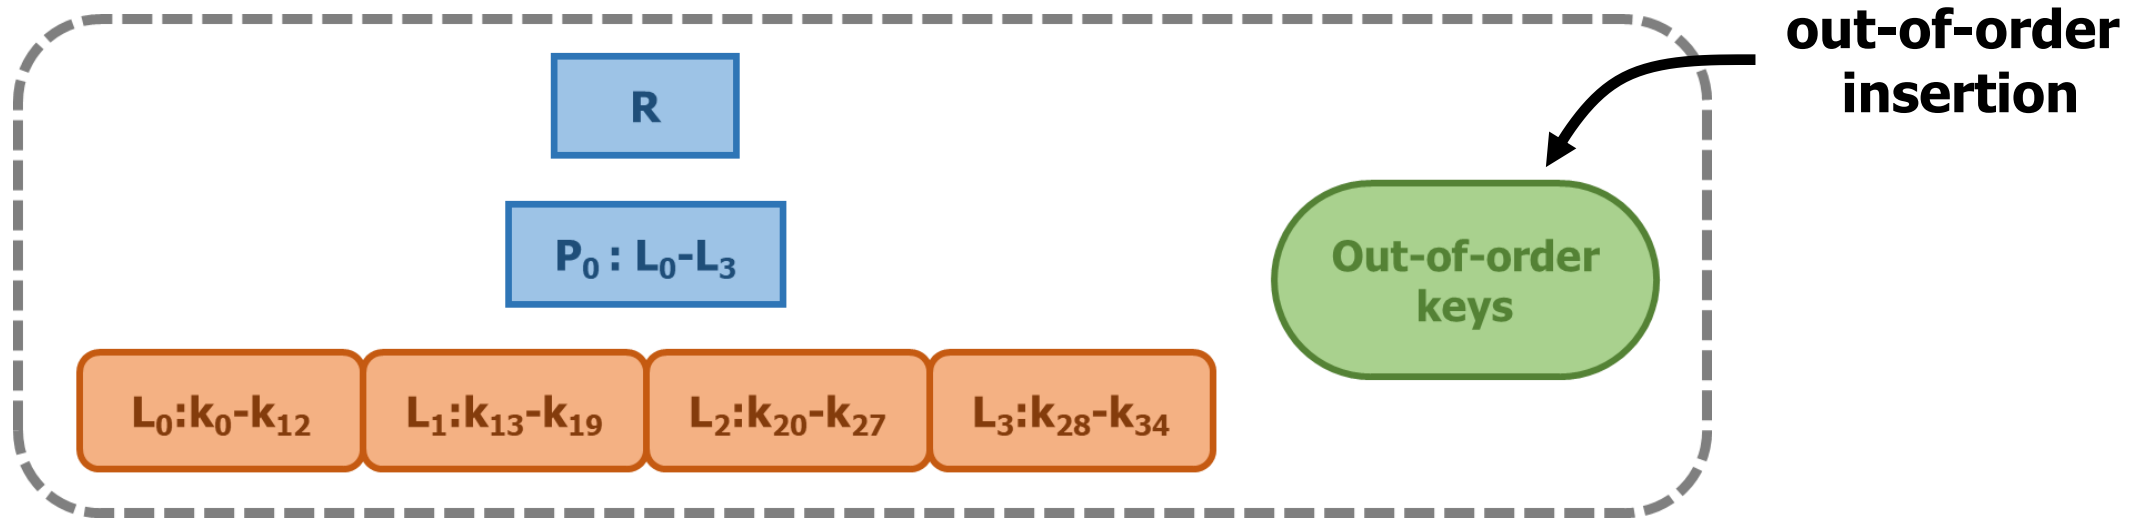

i_piece 

l_piece 



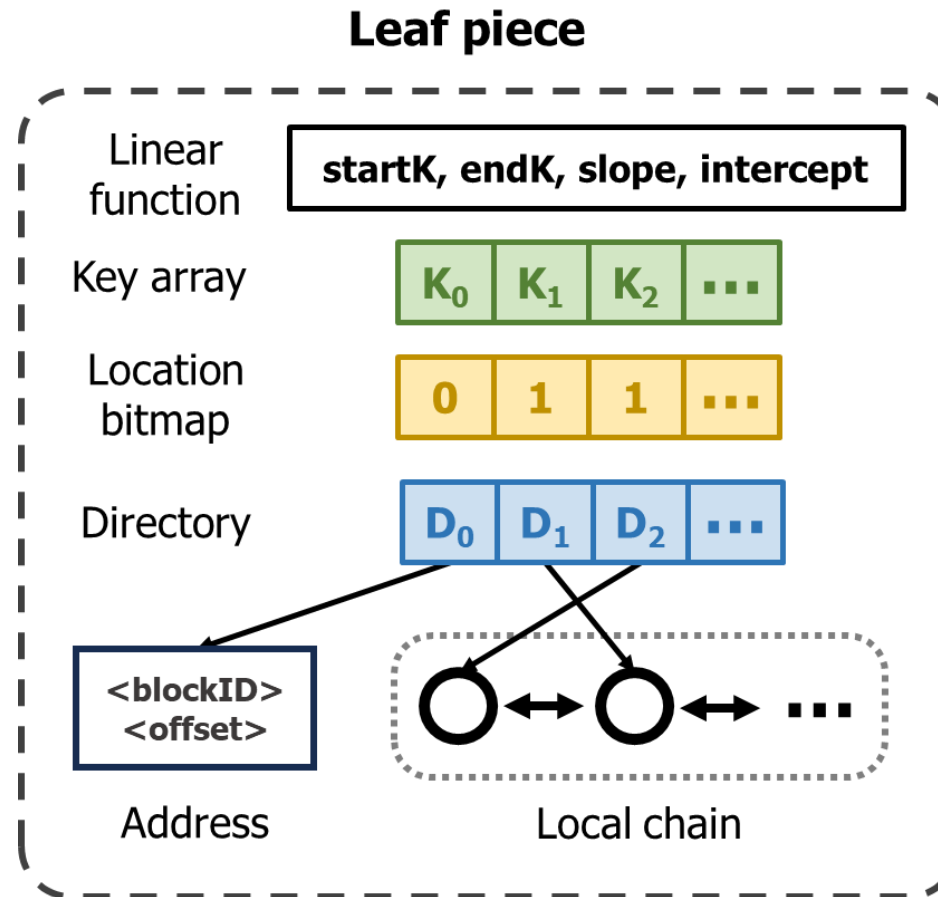
Dynamic Learned Model

- Handling out-of-order insertions
 - Create sort_piece, which is special type of I_piece
 - Directly locates keys using binary search
 - Create a new I_piece when the size of sort_piece reaches the threshold



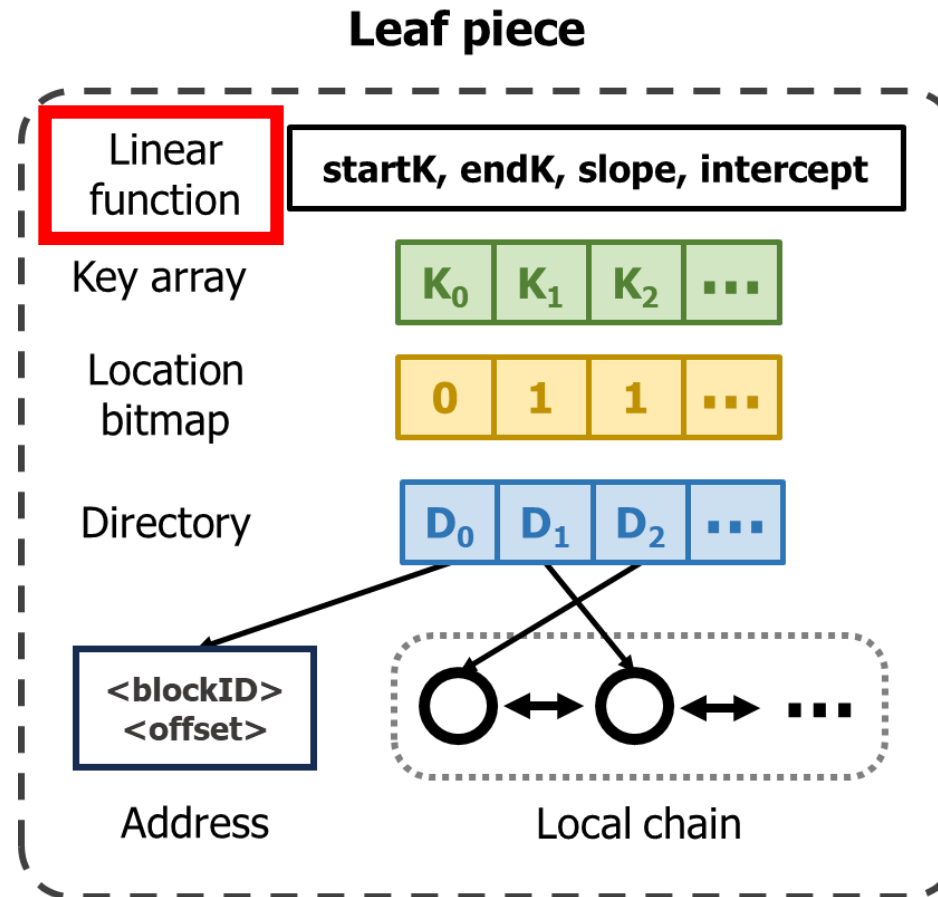
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Location bitmap
 - Directory
 - Local chain
 - Address



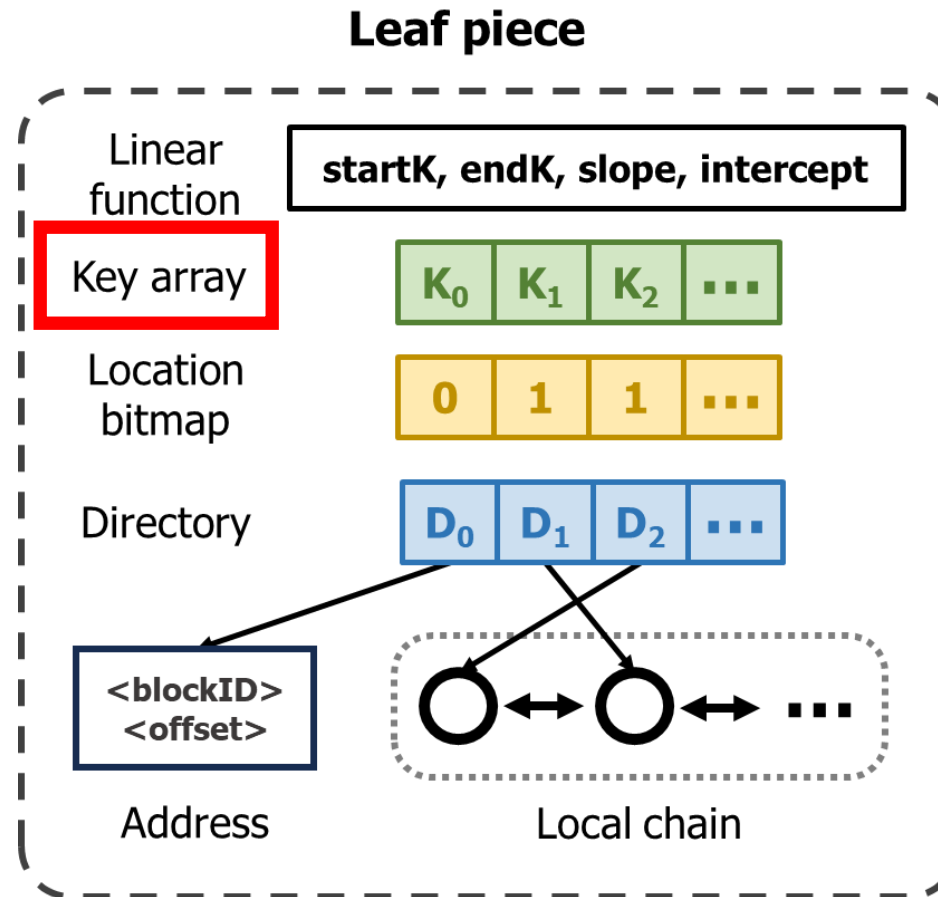
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Predicting position of key
 - Key array
 - Location bitmap
 - Directory
 - Local chain
 - Address



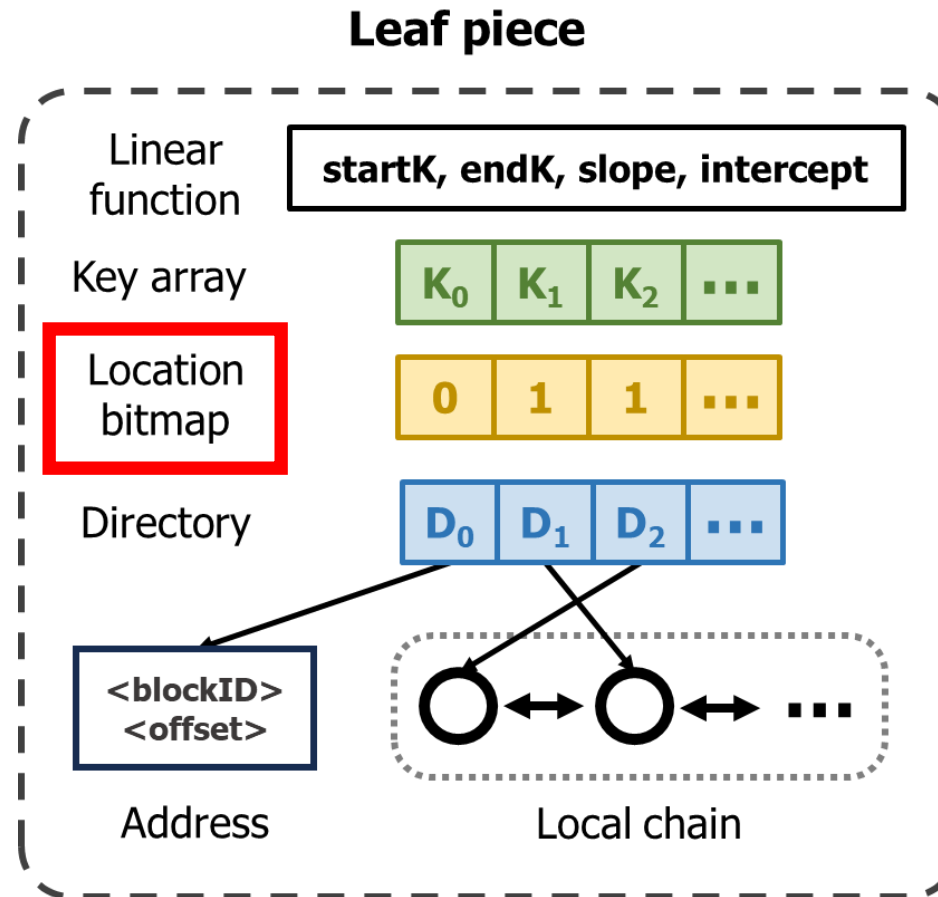
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Recording the keys belonging to leaf piece
 - Location bitmap
 - Directory
 - Local chain
 - Address



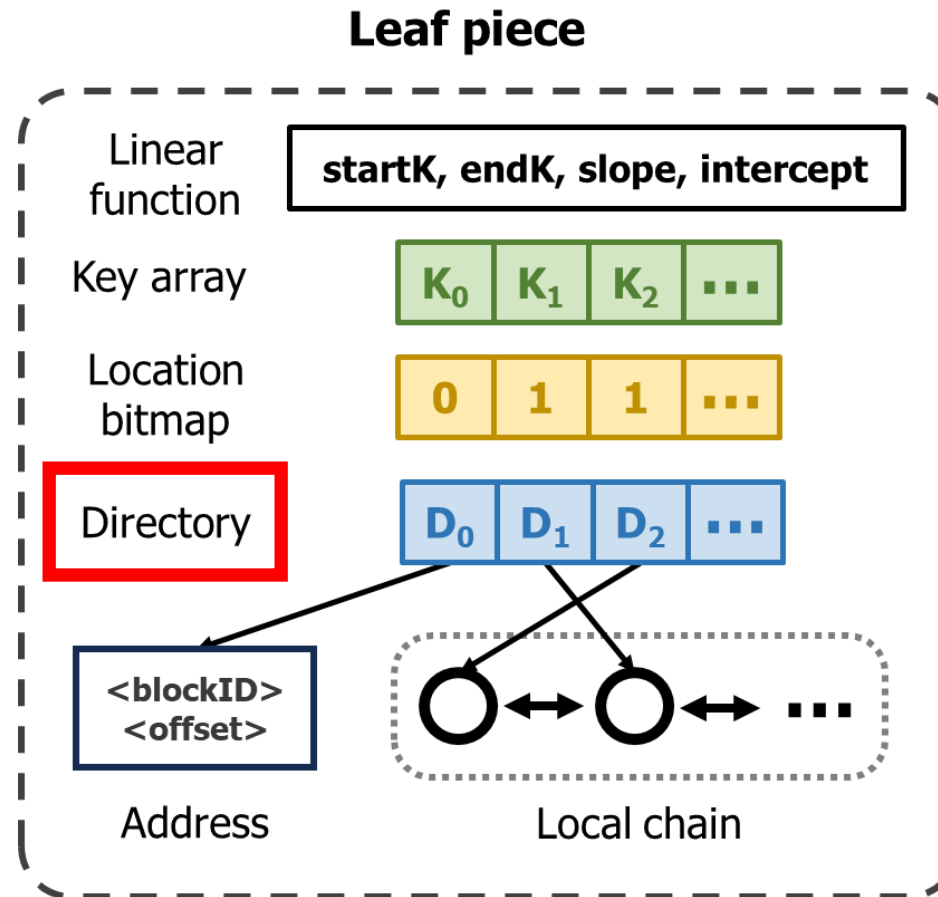
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Location bitmap
 - Tracking whether corresponding data record is in memory or not
 - Directory
 - Local chain
 - Address



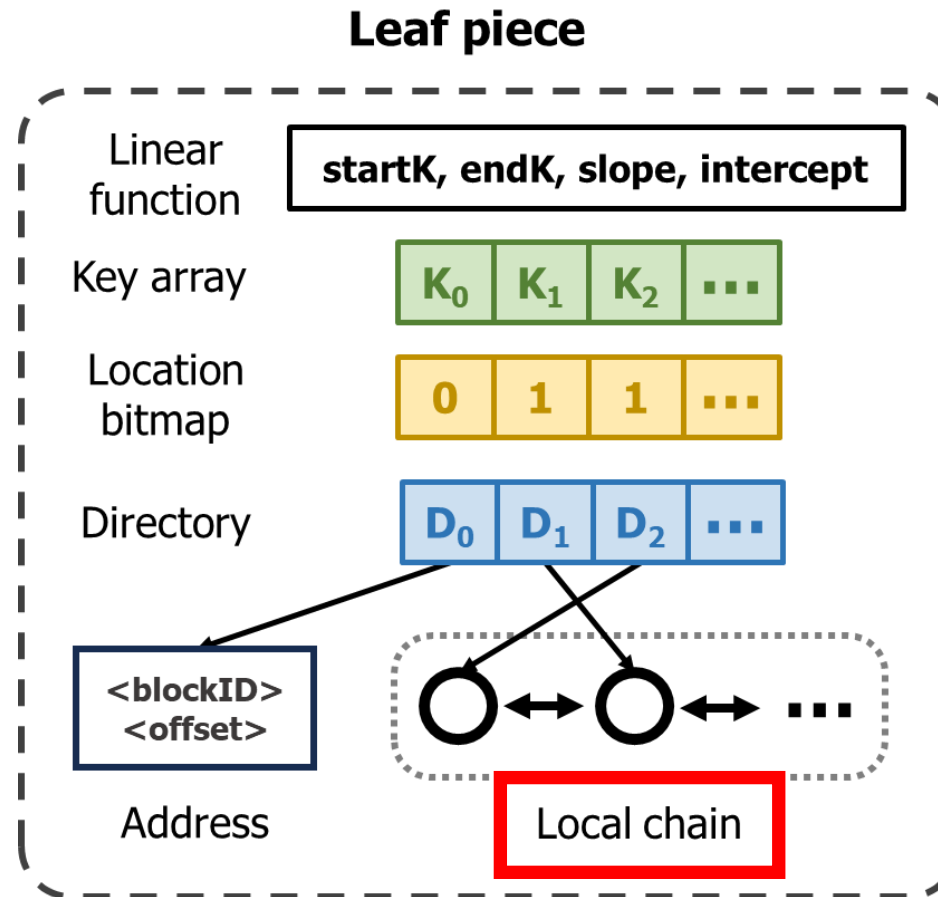
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Location bitmap
 - Directory
 - Tracking the actual storage position of all keys
 - Local chain
 - Address



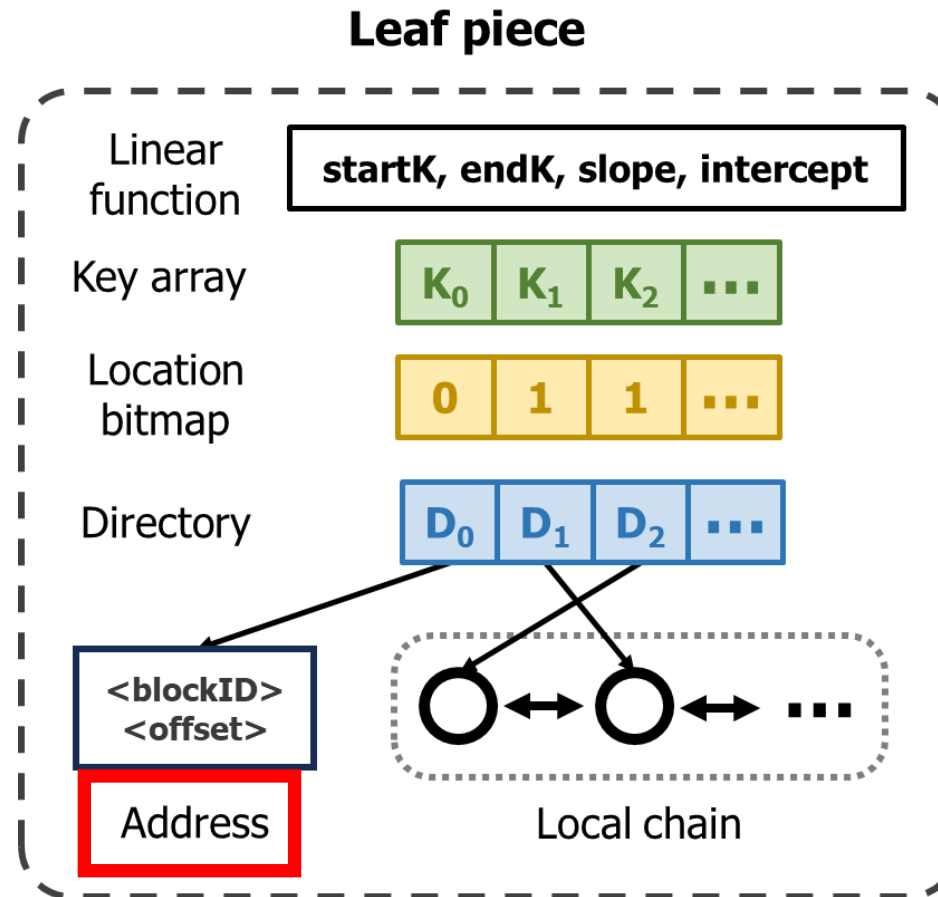
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Location bitmap
 - Directory
 - Stores the payloads of the keys
 - More details in next section
 - Local chain
 - Address



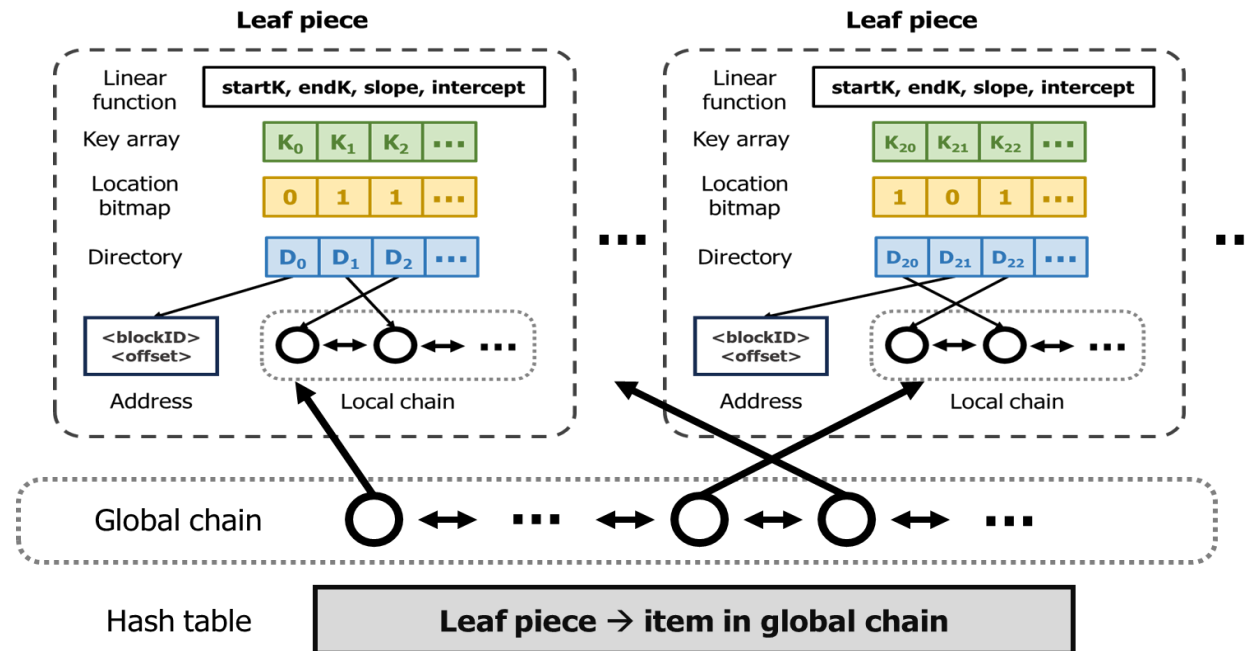
Dynamic Learned Model

- Leaf piece of FILM
 - Linear function
 - Key array
 - Location bitmap
 - Directory
 - Local chain
 - Address
 - Stores the addresses of all evicted records



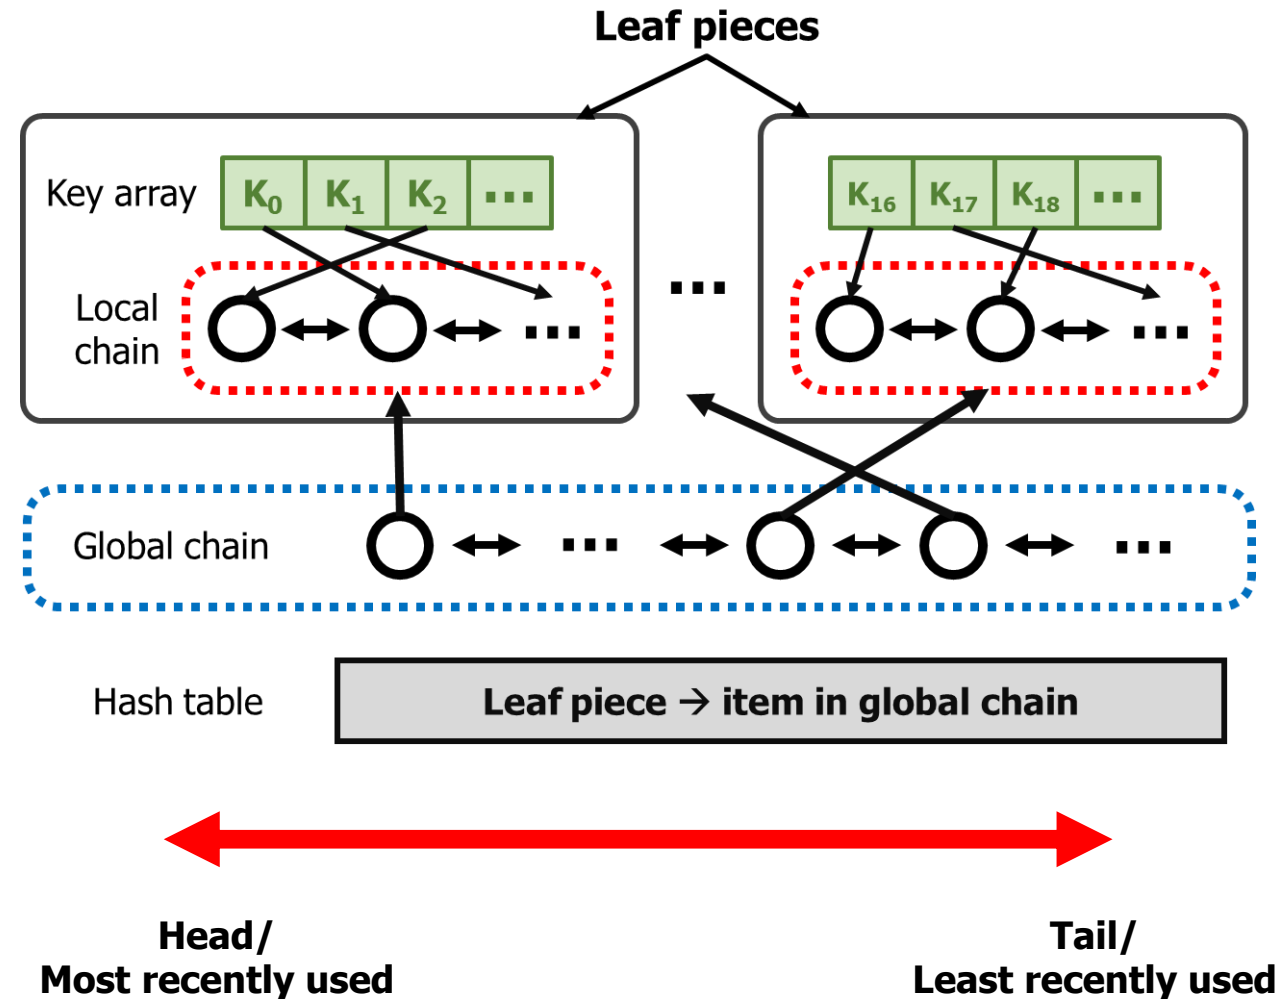
Dynamic Learned Model

- Leaf piece and adaptive LRU of FILM
 - FILM consists of local, global chain to support adaptive LRU
 - Reduce the overhead of maintaining the LRU chain by **piggybacking** the maintenance of the LRU order on the index lookup



Adaptive LRU

- Local chain
 - A component in a leaf piece
 - Maintain LRU order of the keys
 - Stores the payloads of the keys
- Global chain
 - Points to a leaf piece
 - Maintain LRU order of the l_pieces
 - Track the global data access order across pieces
- Hash table
 - For fast search of nodes in global chain



Adaptive LRU

■ Local chain

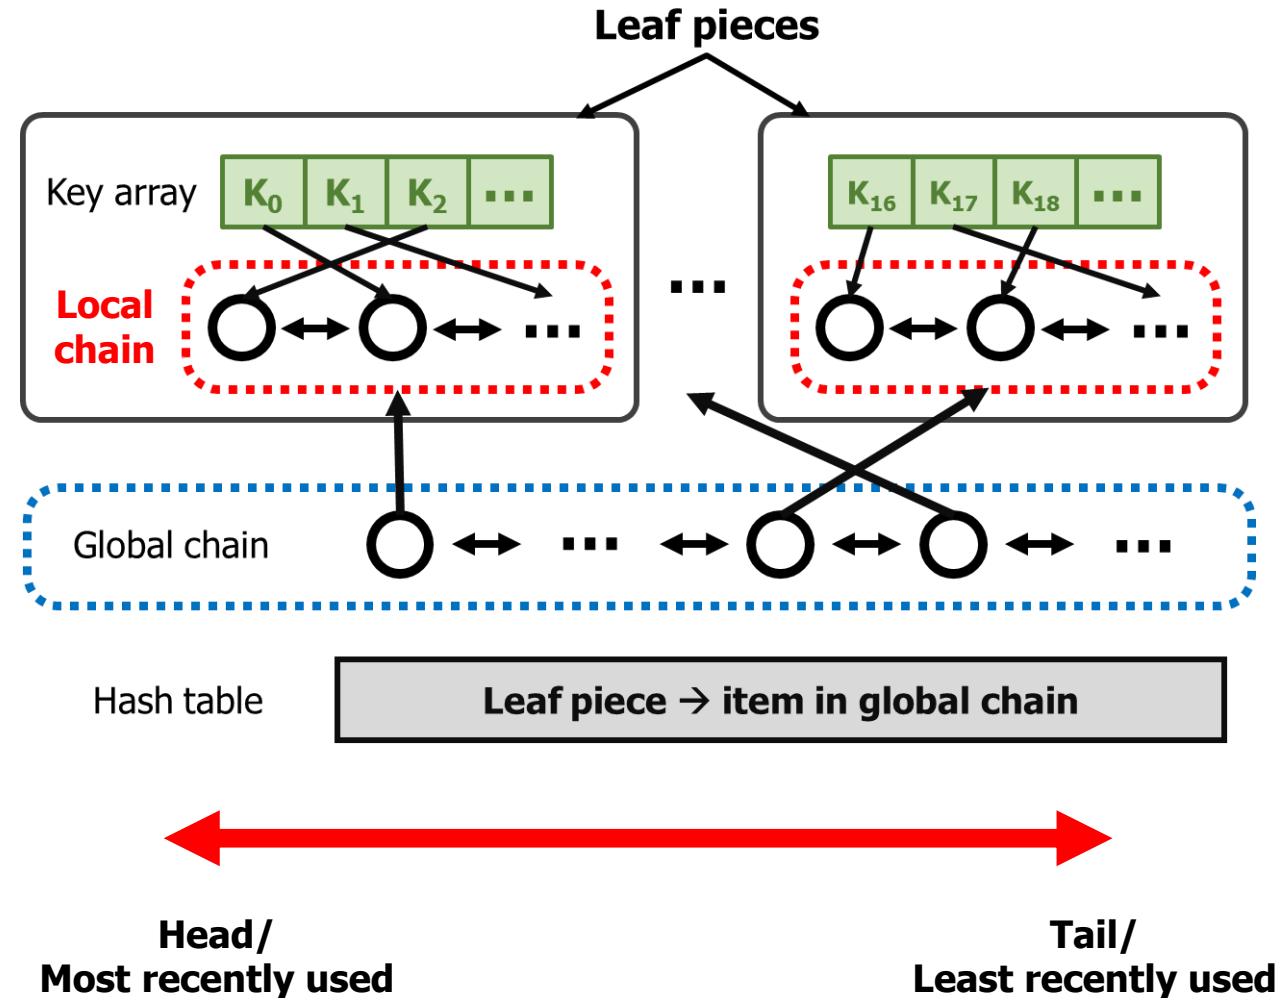
- A component in a leaf piece
- Maintain LRU order of the keys
- Stores the payloads of the keys

■ Global chain

- Points to a leaf piece
- Maintain LRU order of the l_pieces
- Track the global data access order across pieces

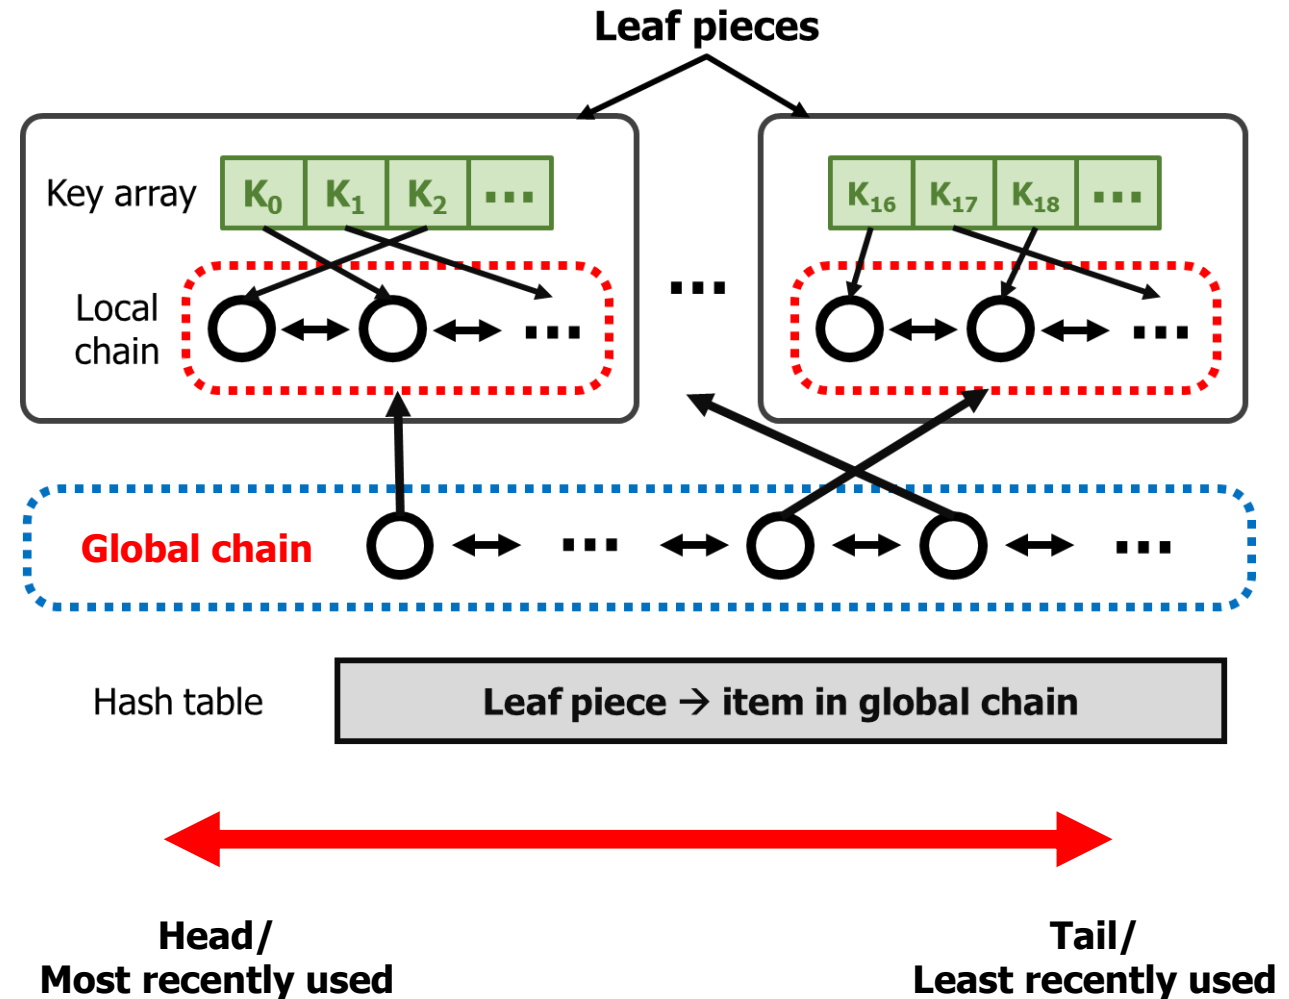
■ Hash table

- For fast search of nodes in global chain



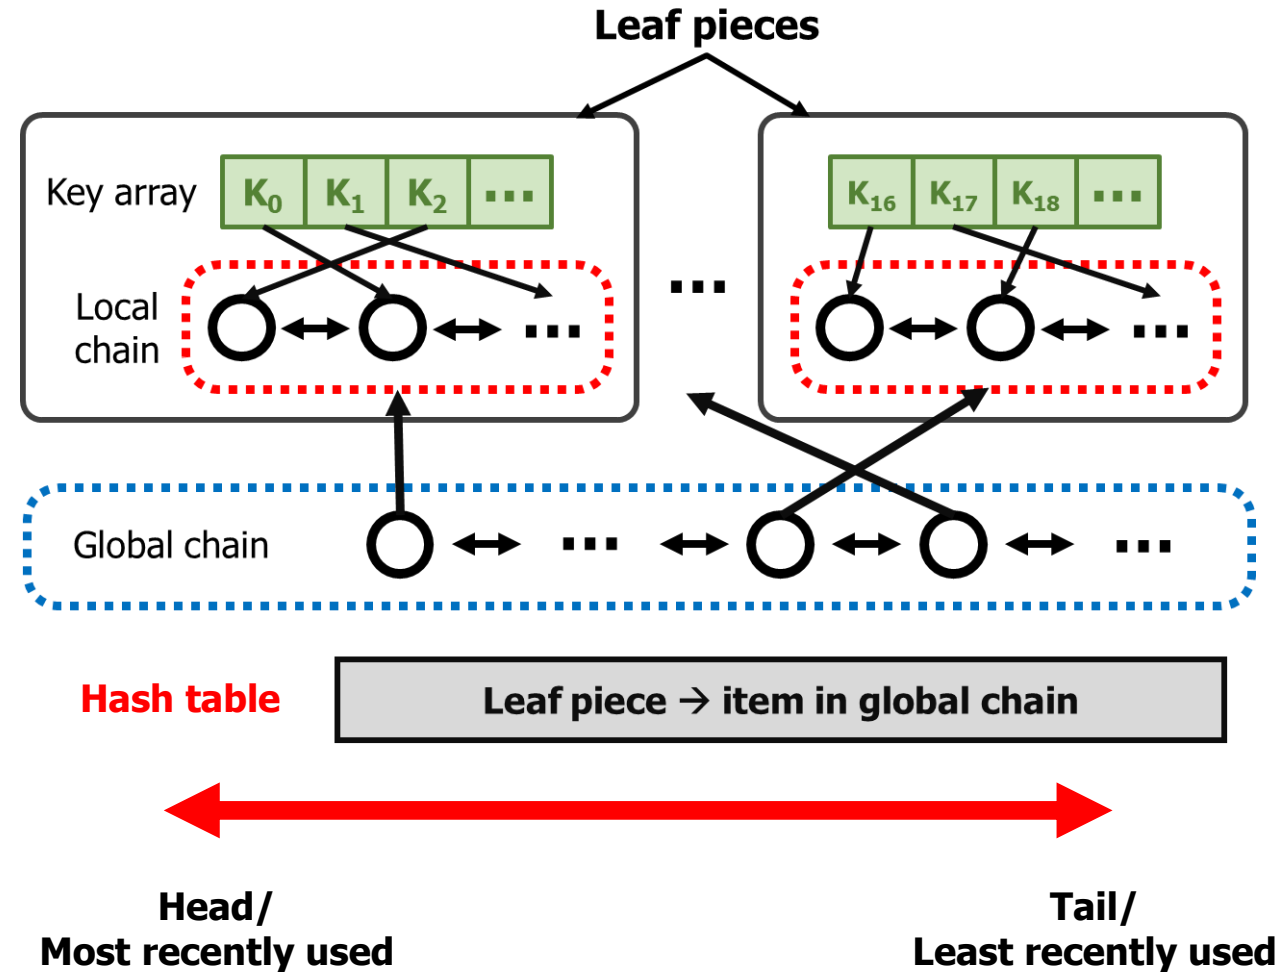
Adaptive LRU

- Local chain
 - A component in a leaf piece
 - Maintain LRU order of the keys
 - Stores the payloads of the keys
- Global chain
 - Points to a leaf piece
 - Maintain LRU order of the l_pieces
 - Track the global data access order across pieces
- Hash table
 - For fast search of nodes in global chain



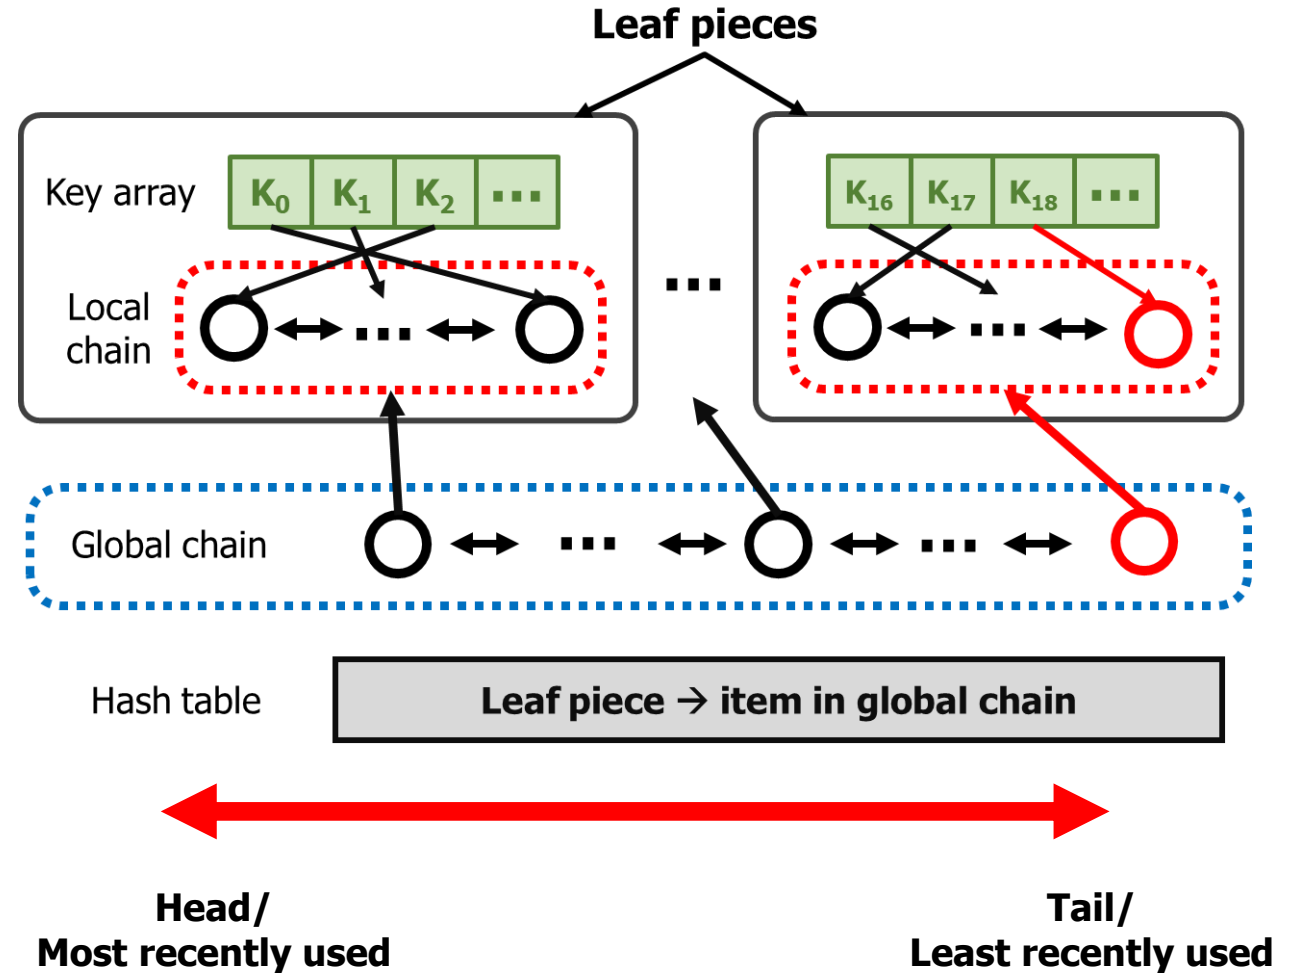
Adaptive LRU

- Local chain
 - A component in a leaf piece
 - Maintain LRU order of the keys
 - Stores the payloads of the keys
- Global chain
 - Points to a leaf piece
 - Maintain LRU order of the l_pieces
 - Track the global data access order across pieces
- Hash table
 - For fast search of nodes in global chain



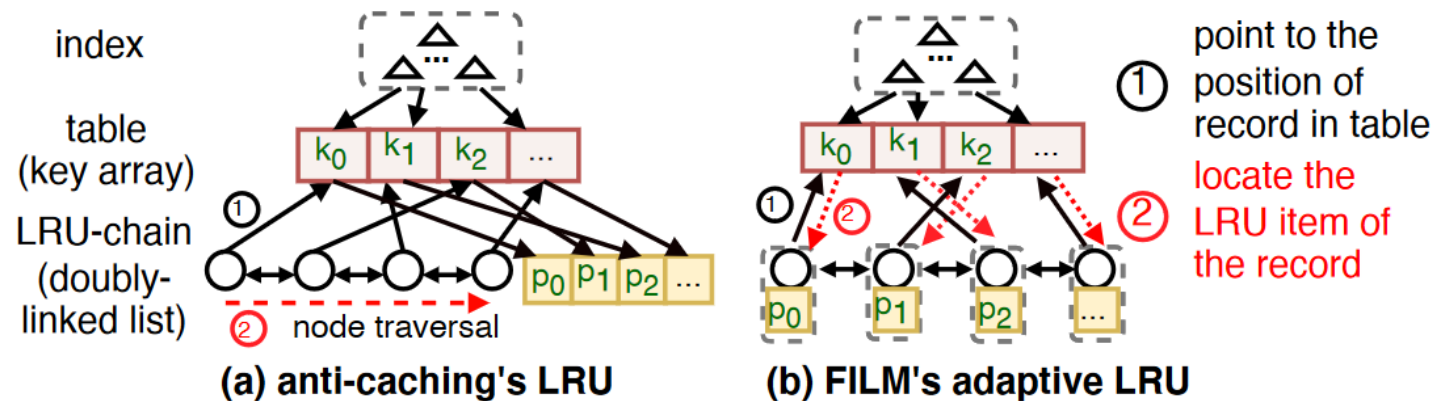
Adaptive LRU

- Cold Data identification
 - When available memory runs out
 - Global chain identifies cold leaf piece to eviction
 - Evicts records from the coldest leaf to a block until block is full
 - Continue until database size reaches a user-specified threshold



Adaptive LRU

- Difference between FILM's adaptive LRU and anti-caching's LRU
 - FILM avoids the node traversal by piggybacking the locating of the LRU item onto the index lookup
 - An LRU item in FILM's adaptive LRU stores the payload of the corresponding record



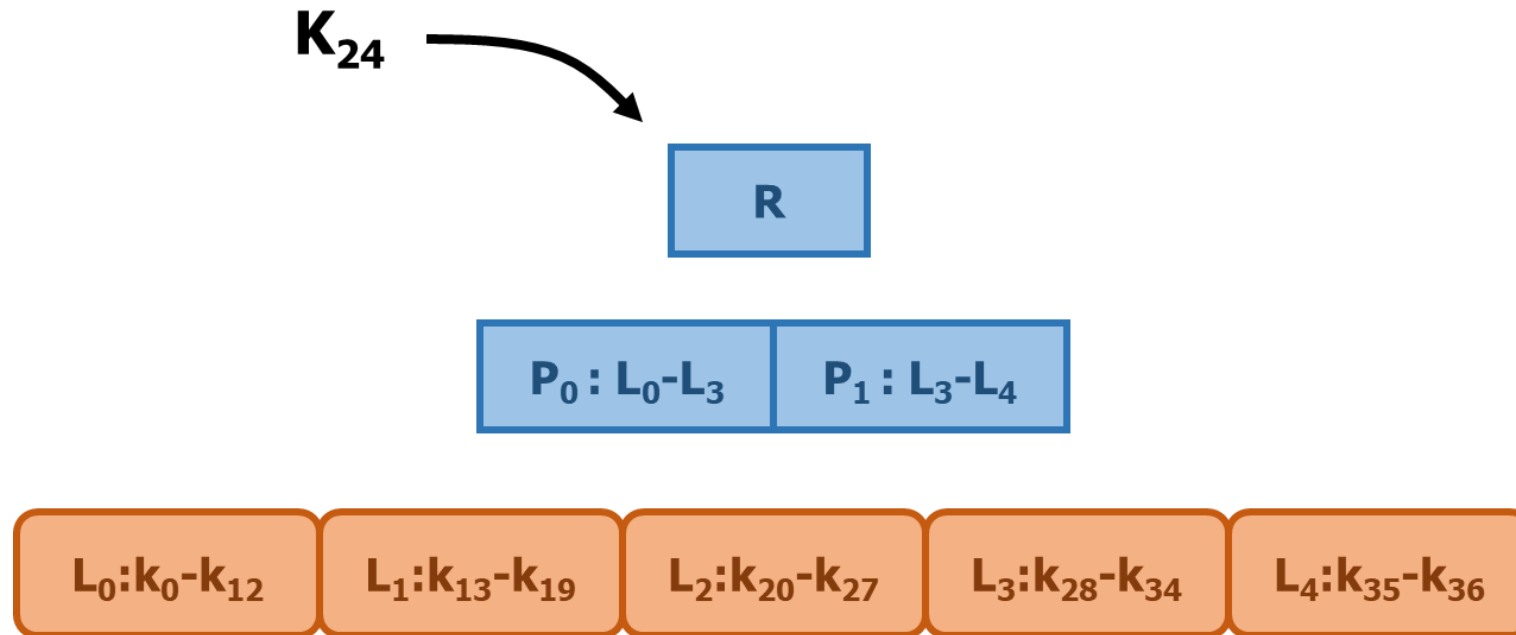
Note: In anti-caching, the pointers of LRU chain are embeds in the records' headers. To show the difference when locating LRU items more intuitive, we separate them in this toy example.

Fig. 4: LRU in anti-caching vs. adaptive LRU in FILM

Query Processing

- Point Query

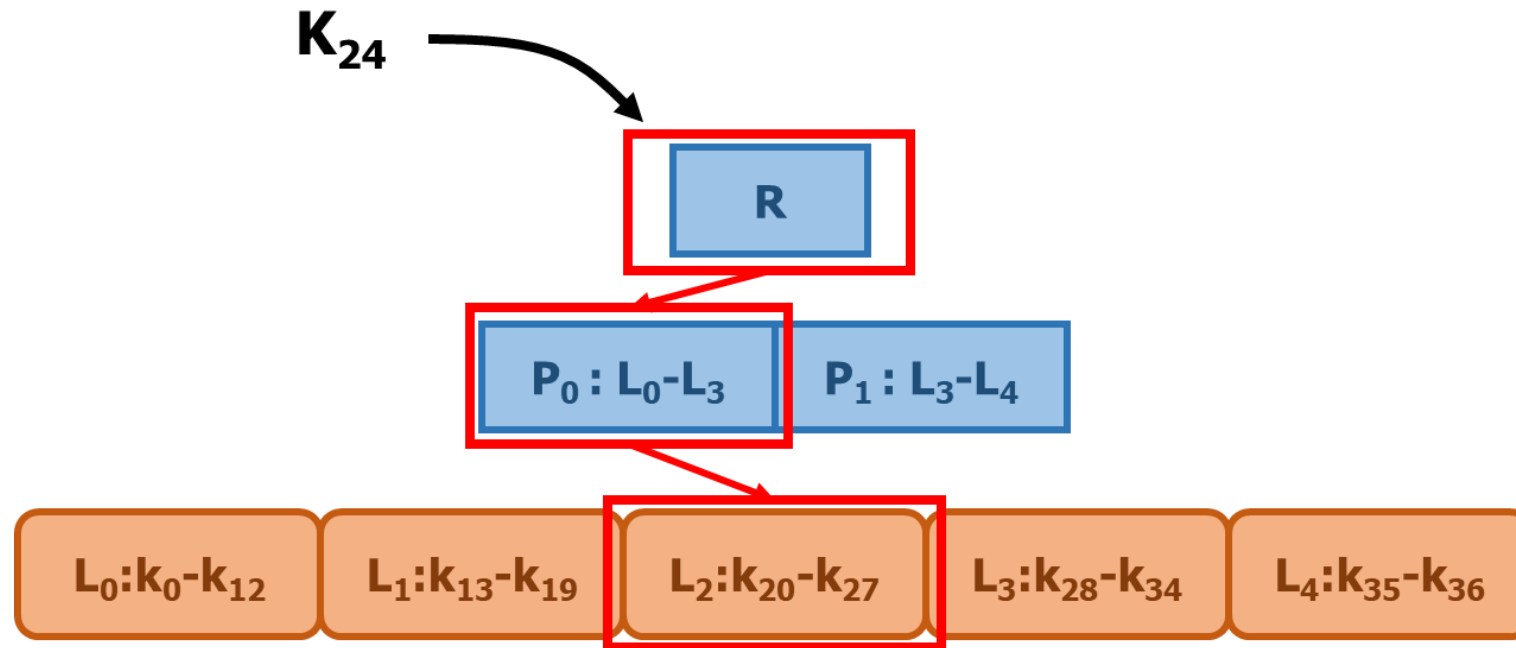
- Find the leaf piece that **K** belongs to
Recursively predict the candidate pieces



Query Processing

- Point Query

- Find the leaf piece that **K** belongs to
Recursively predict the candidate pieces

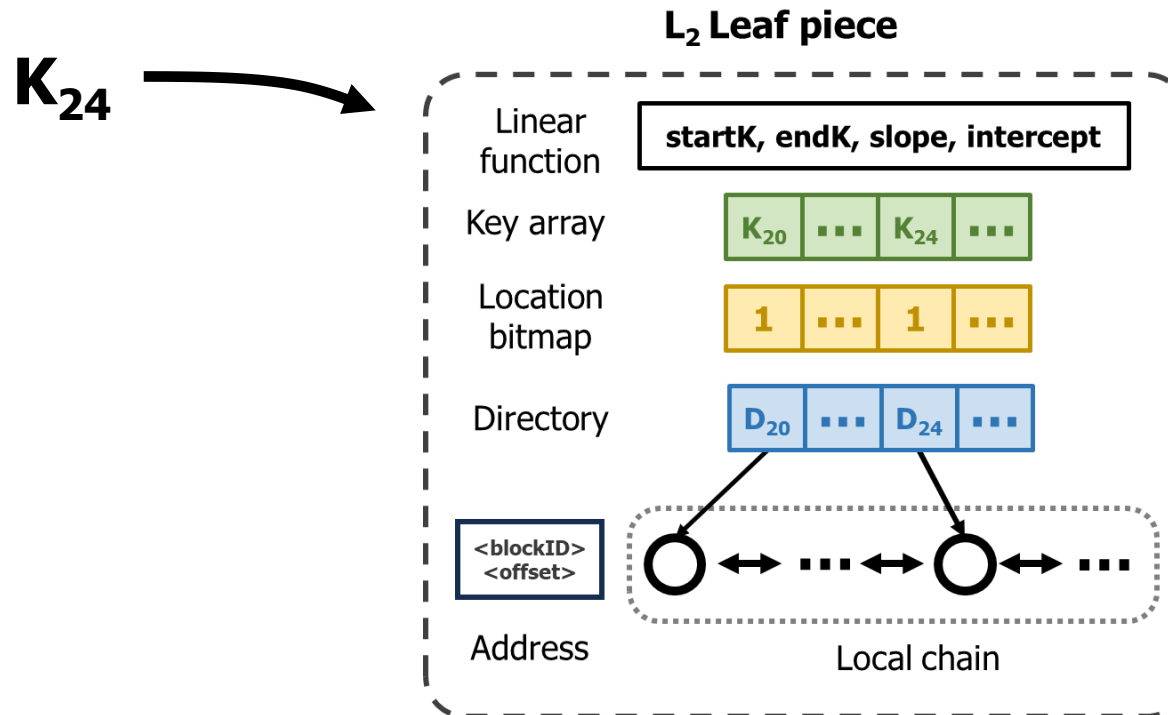


Query Processing

- Point Query

- 2. Locate **K** in the piece

- Predict the position of **K**

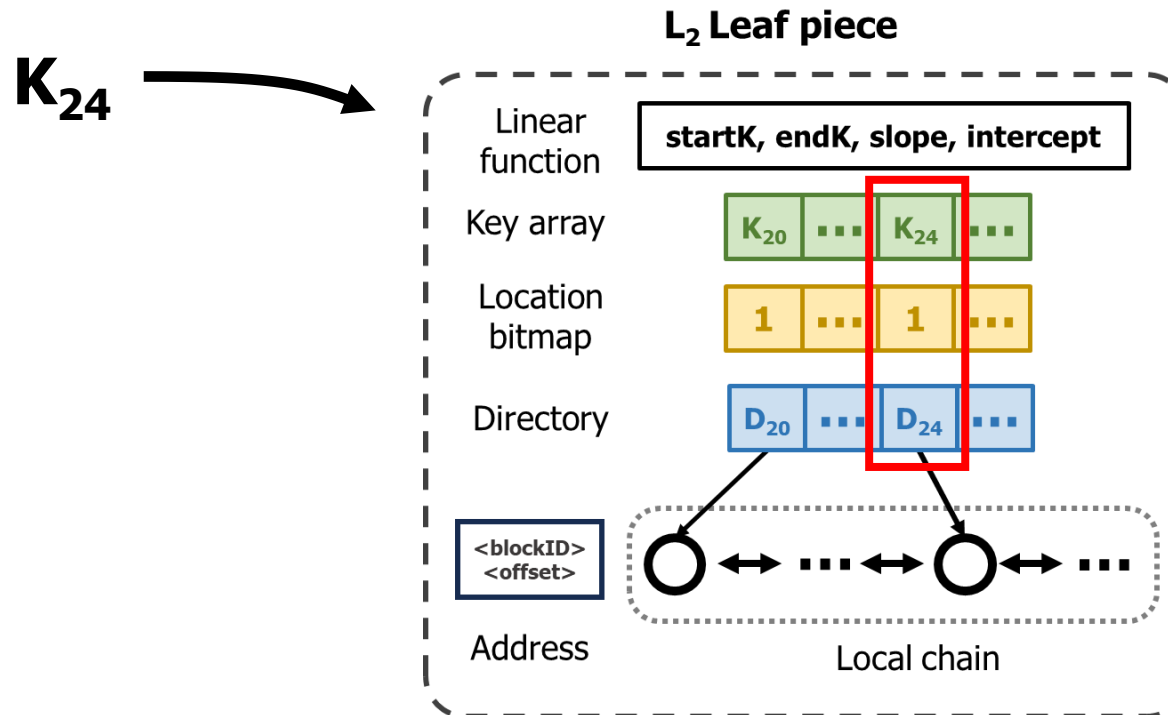


Query Processing

- Point Query

- 2. Locate **K** in the piece

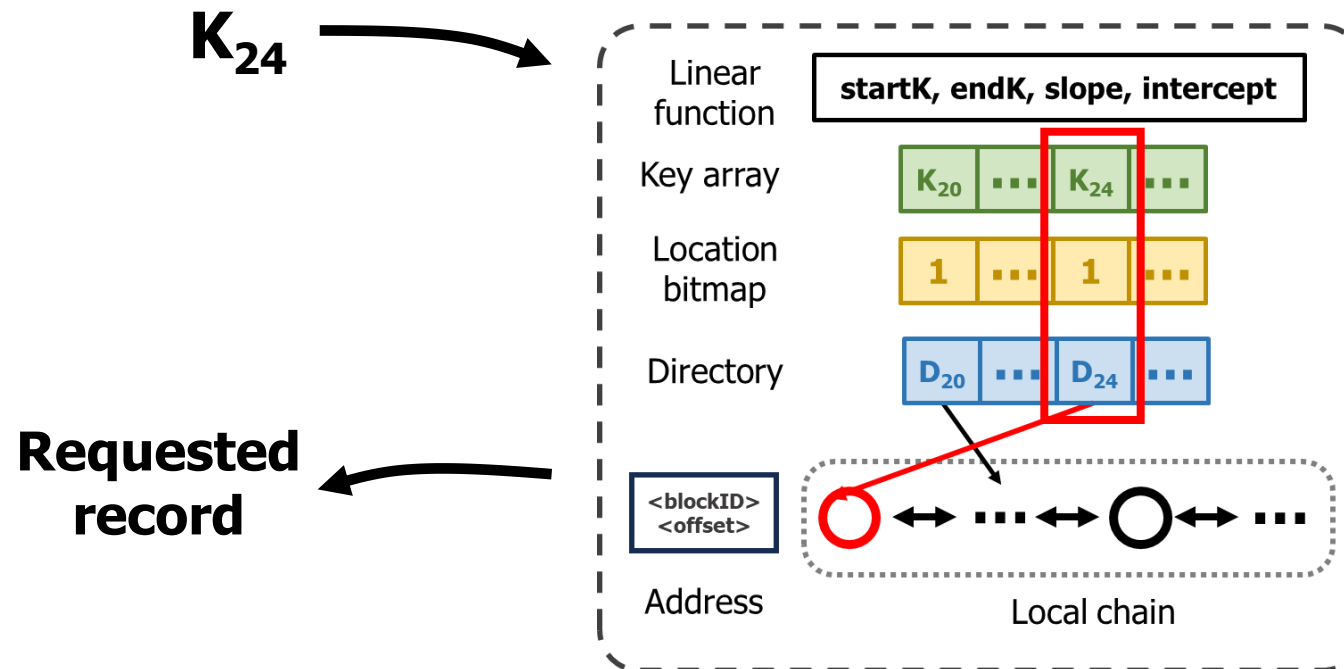
FILM uses the location bitmap to know if k's payload is in memory or on disk



Query Processing

- Point Query

- Retrieve the data record and update the adaptive LRU
Record can be directly accessed using local chain or address

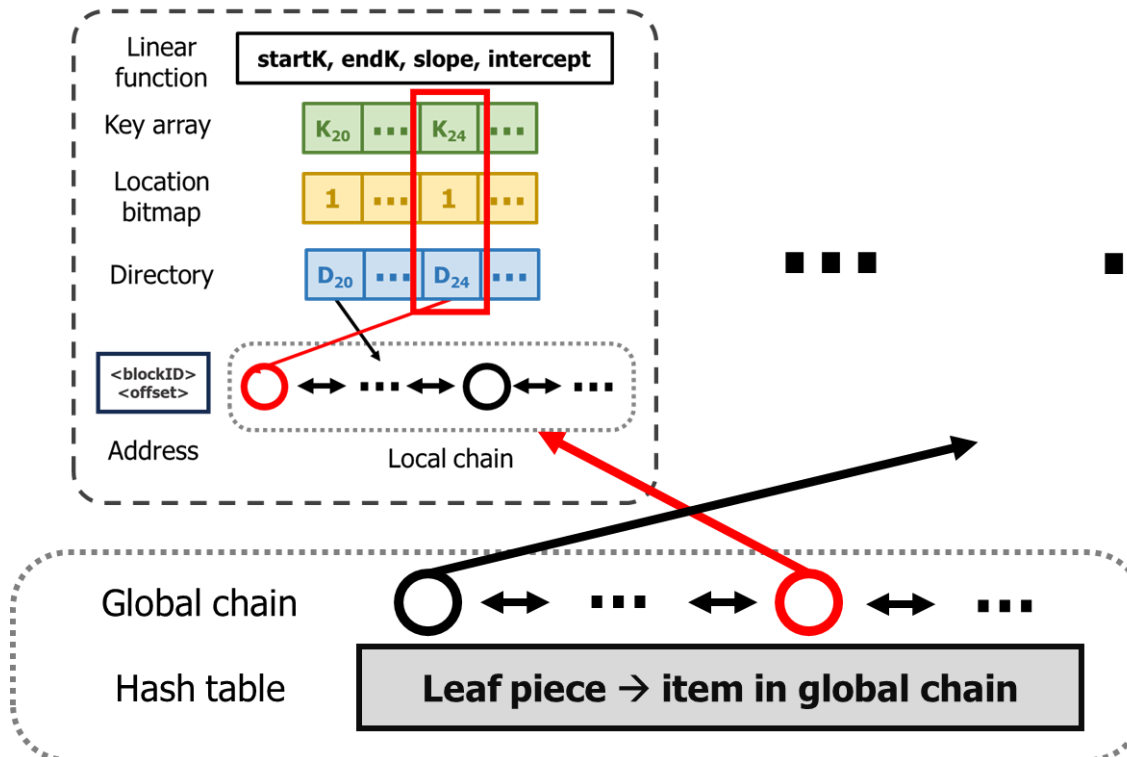


Query Processing

- Point Query

3. Retrieve the data record and update the adaptive LRU

The accessed leaf will be moved to the head of the global chain

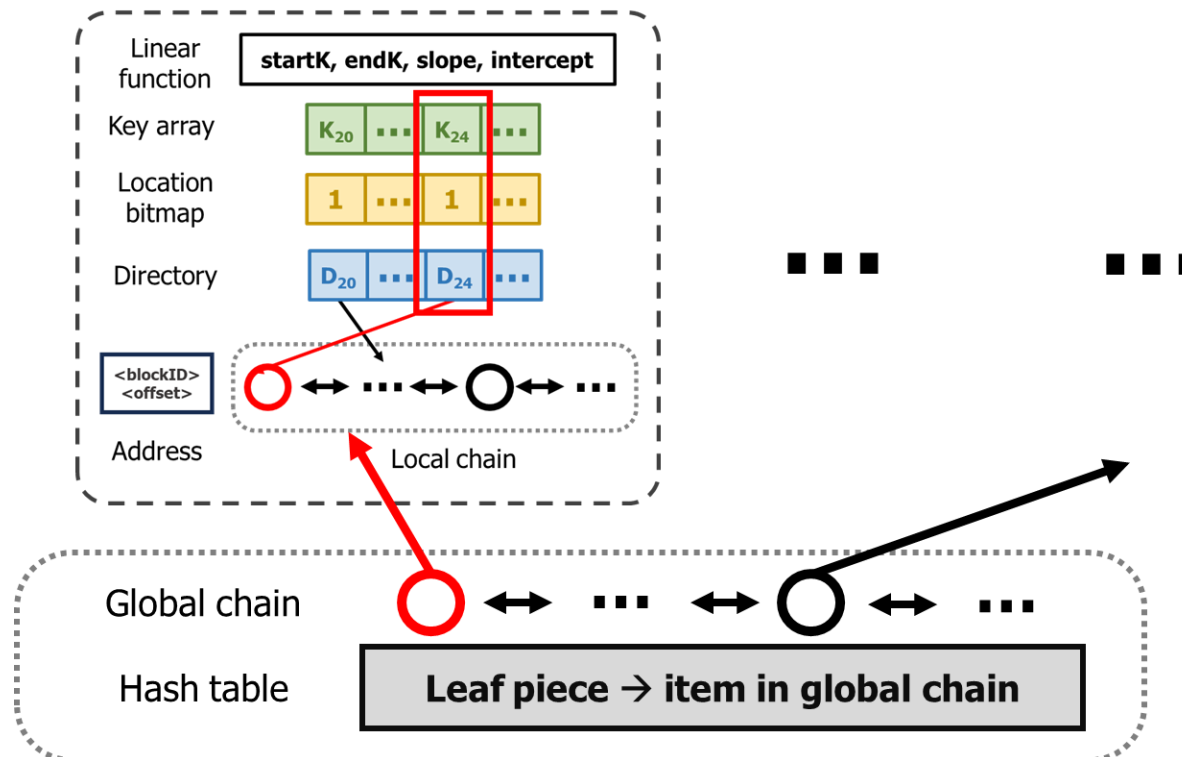


Query Processing

- Point Query

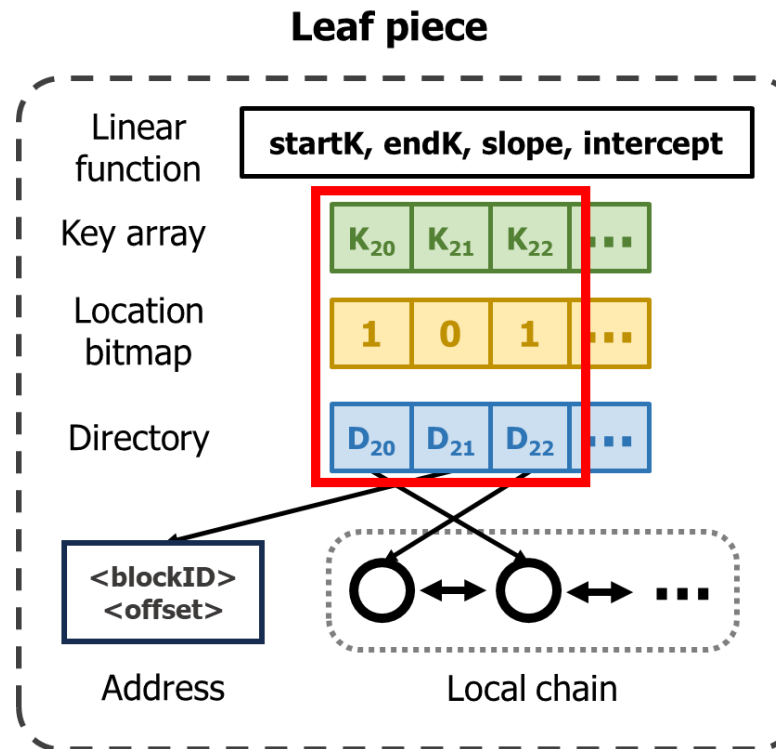
3. Retrieve the data record and update the adaptive LRU

The accessed leaf will be moved to the head of the global chain



Query Processing

- Range Query
 - FILM handles requested data records similarly whether they reside in a single type or multiple types of storage devices.



Experiments

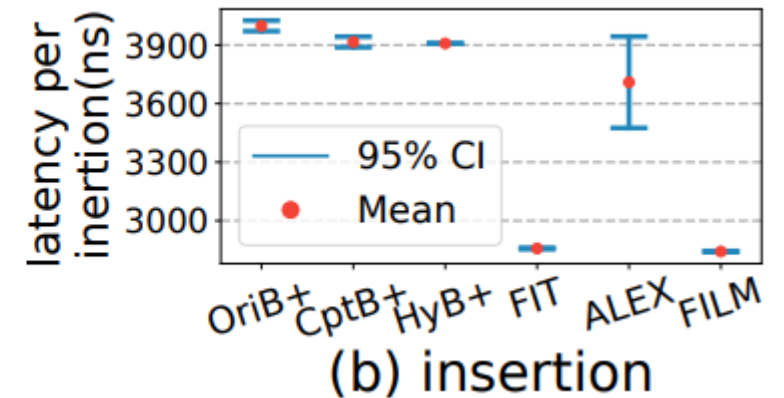
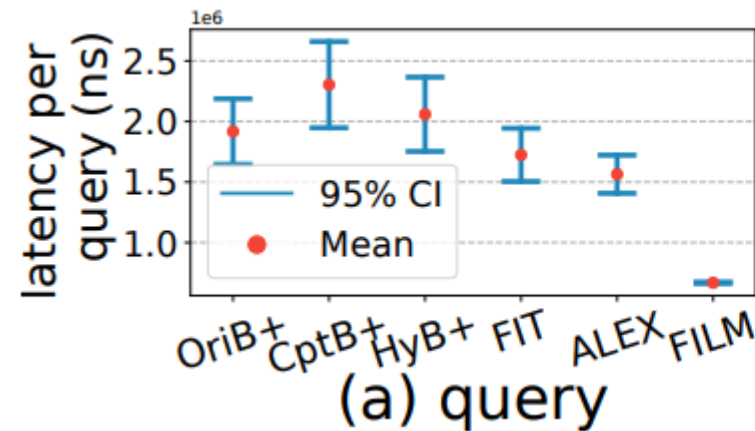
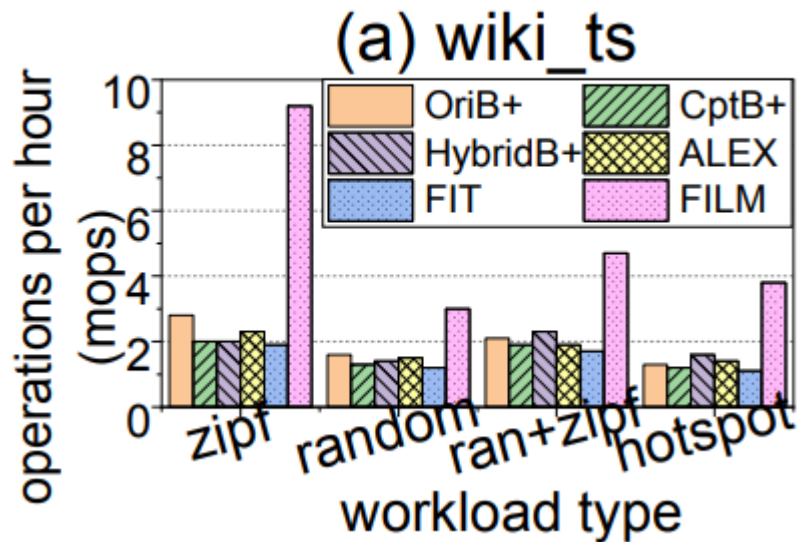
- Baselines and Experiment Setup

- Comparison(Base Line) : B+tree, HybridB+tree, CptB+tree, Alex, Fitting-Tree
- Datasets : wiki_ts, books, astro_a, synthetic, YCSB (64bit Key, 128Byte Value)
- Workload : Zipfian, Hotspot, Random, Zipfian+Random
- Spec :
 - 3.6GHz Intel CPU with 256KB L1 cache
 - 128GB memory (4 × 32GB)
 - 557GB disk
 - Single thread
 - Direct I/O

Experiments

■ Comparison with Baselines

- Insertion and query performance(Avg. 4GB All dataset, 1:1 , 1 hour)

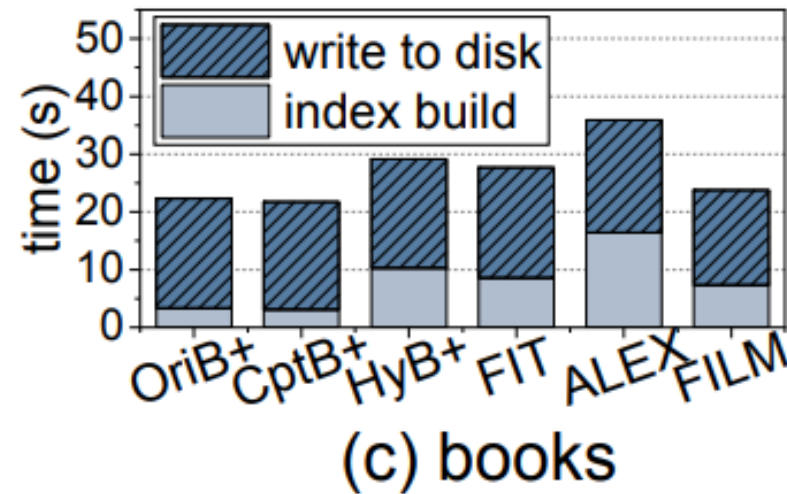


FILM has a **2-5 times advantage** over base line in terms of insertion and query performance.

Film also has the **lowest latency and more stable performance**.

Experiments

- Comparison with Baselines
 - Index construction



The learned index has **advantages** over traditional indexes **in memory usage**, but correspondingly, it has **disadvantages in build time** (about 2.2~2.5 times).

Also, due to occupying less memory, it has more **advantages in writing to disk**.


Experiments

■ Comparison with Baselines


- LRU overhead(books, 4GB dataset, 100000 random queries)

aLRU : LRU with a sampling rate of 0.01

adaLRU : Adaptive LRU



(a) aLRU vs. adaptive LRU





(b) adaB+tree vs. FILM

On LRU, adaLRU is superior to aLRU.
On the index, FILM is superior to traditional indexes.

Experiments

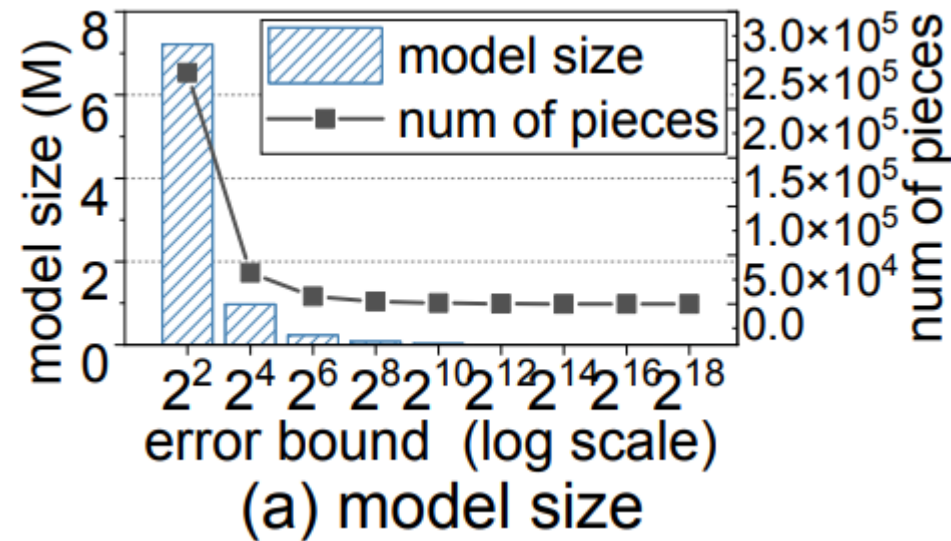
- Study of Environment Parameters
 - Available memory(4GB wiki_ts dataset, Zipfian workload)



Increasing memory can increase the number of operations per unit time and reduce disk io, with FILM gaining more significant benefits in this regard.

Experiments


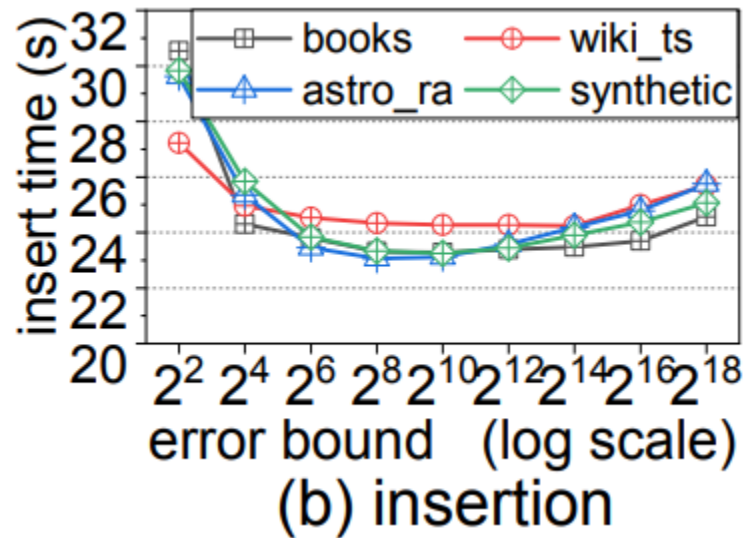
- Sensitivity to Data and Model Parameters
 - Error bound ε



When the error bound increases from 4 to 16, the sharp decrease in the number of pieces fitted by FILM leads to a decrease in the memory occupied by FILM, which then stabilizes.

Experiments


- Sensitivity to Data and Model Parameters
 - Error bound ϵ



FILM within the range of 16 to 2^{12} ϵ having the best performance.

Experiments

- Sensitivity to Data and Model Parameters
 - Handling Out-of-Order Insertions(sort_piece)



Both too small and too large sort pieces can lead to performance loss.
30000~70000 is the optimal range.

Thank you